

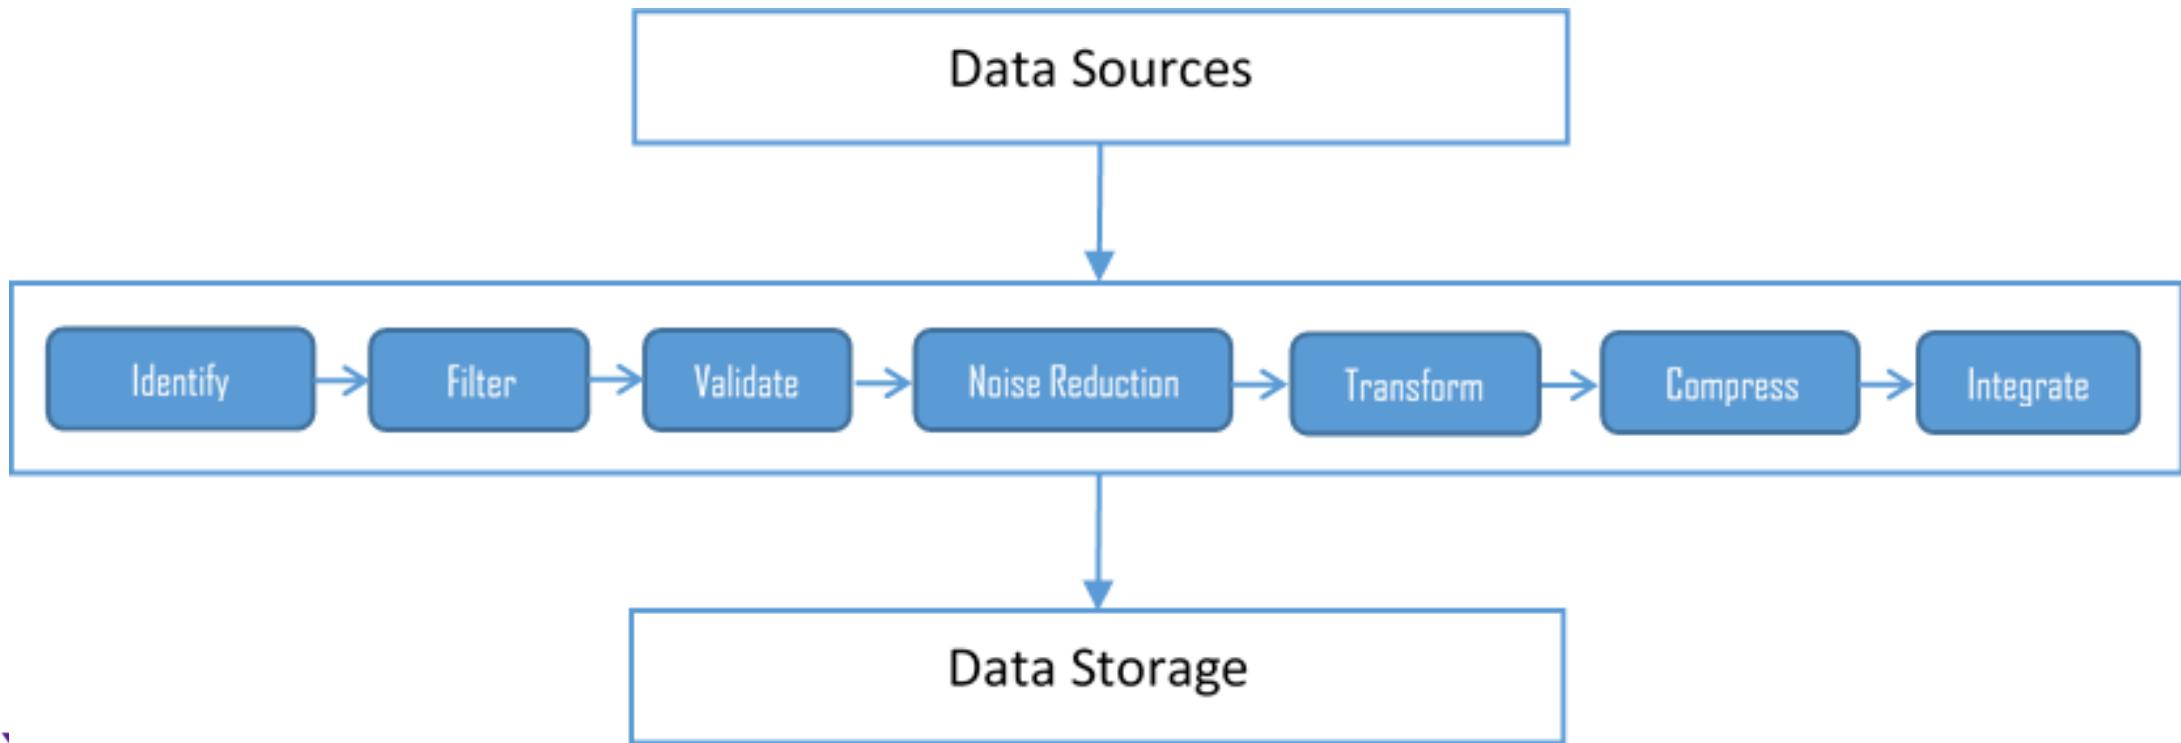
Big Data Patterns

Design patterns for big data

- reduce complexity
- boost performance
- improve results for new and larger forms of data
- reference: <https://www.packtpub.com/application-development/architectural-patterns>

Data sources and ingestion layer

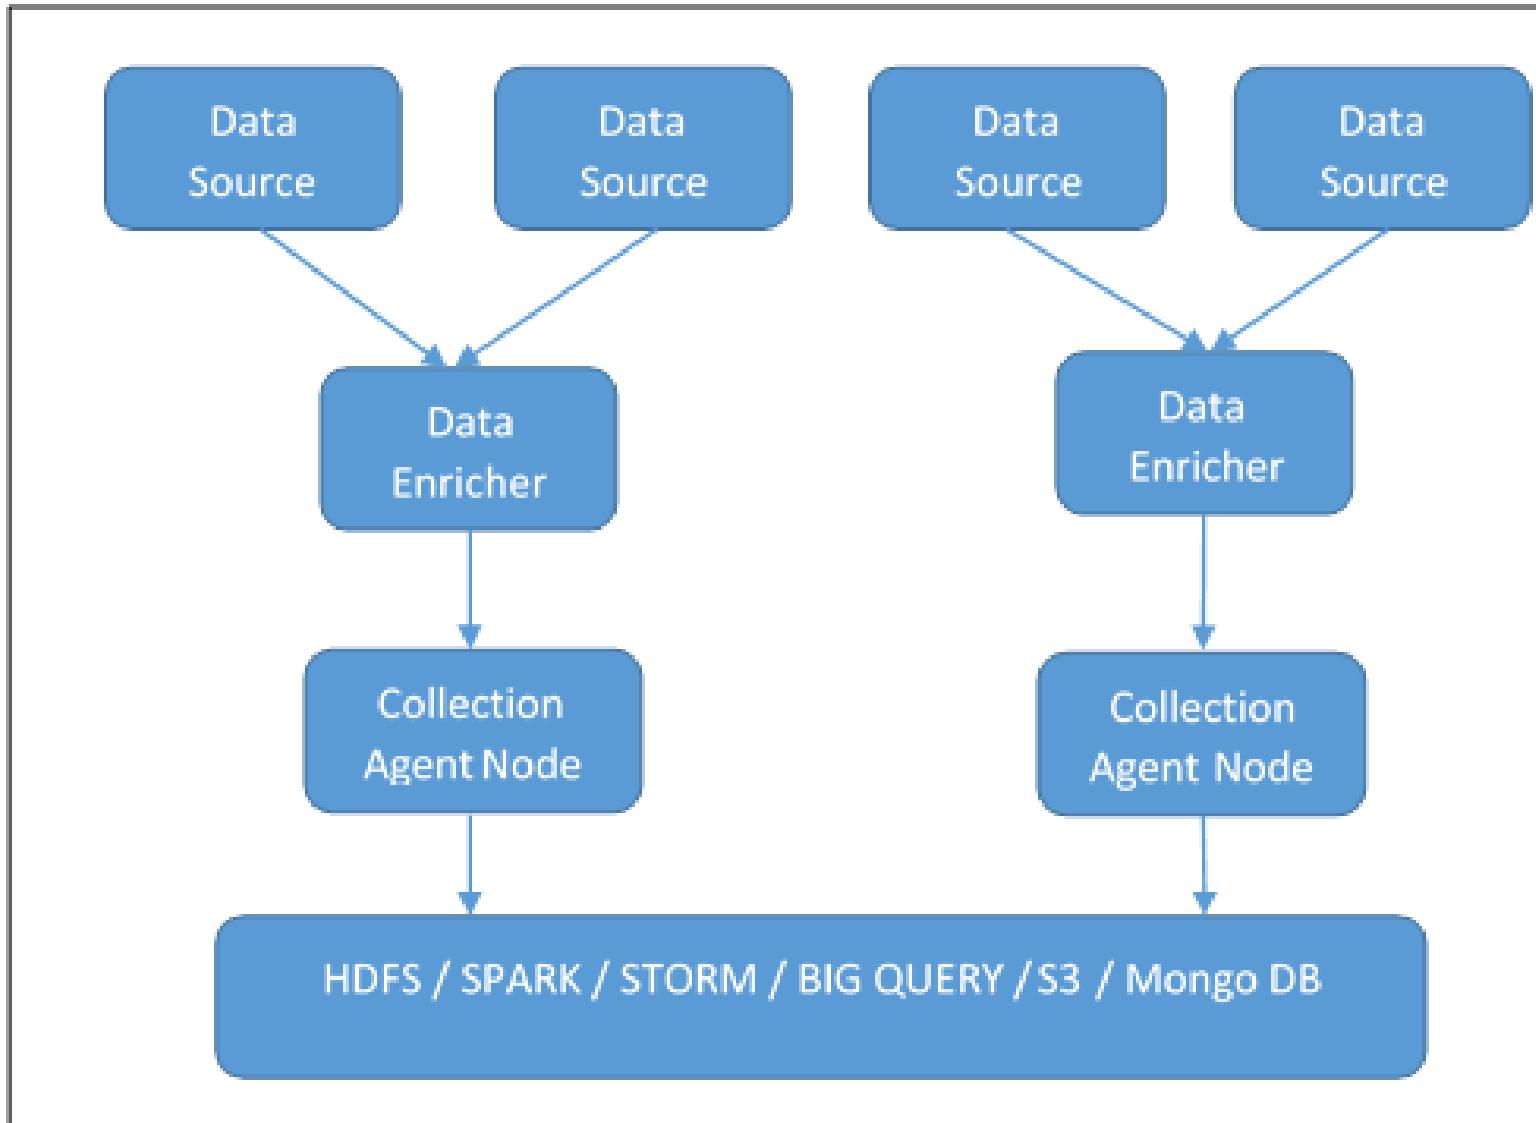
- Multiple data source load and prioritization
- Ingested data indexing and tagging
- Data validation and cleansing
- Data transformation and compression

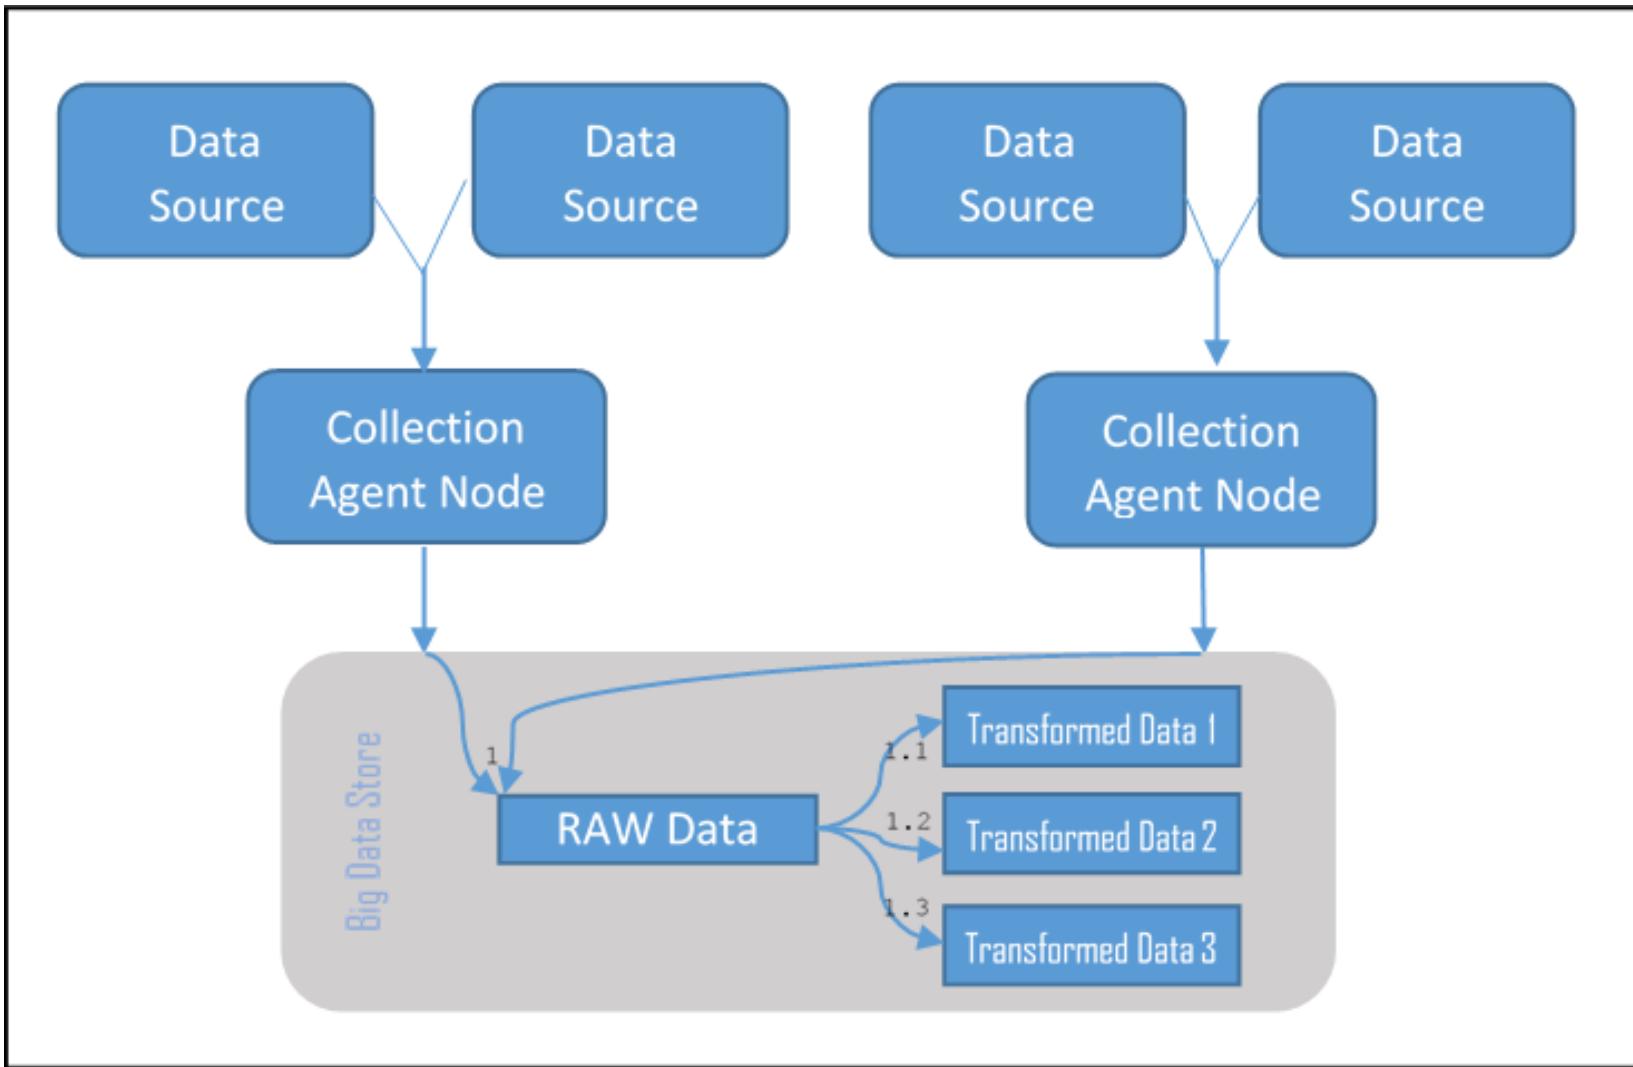


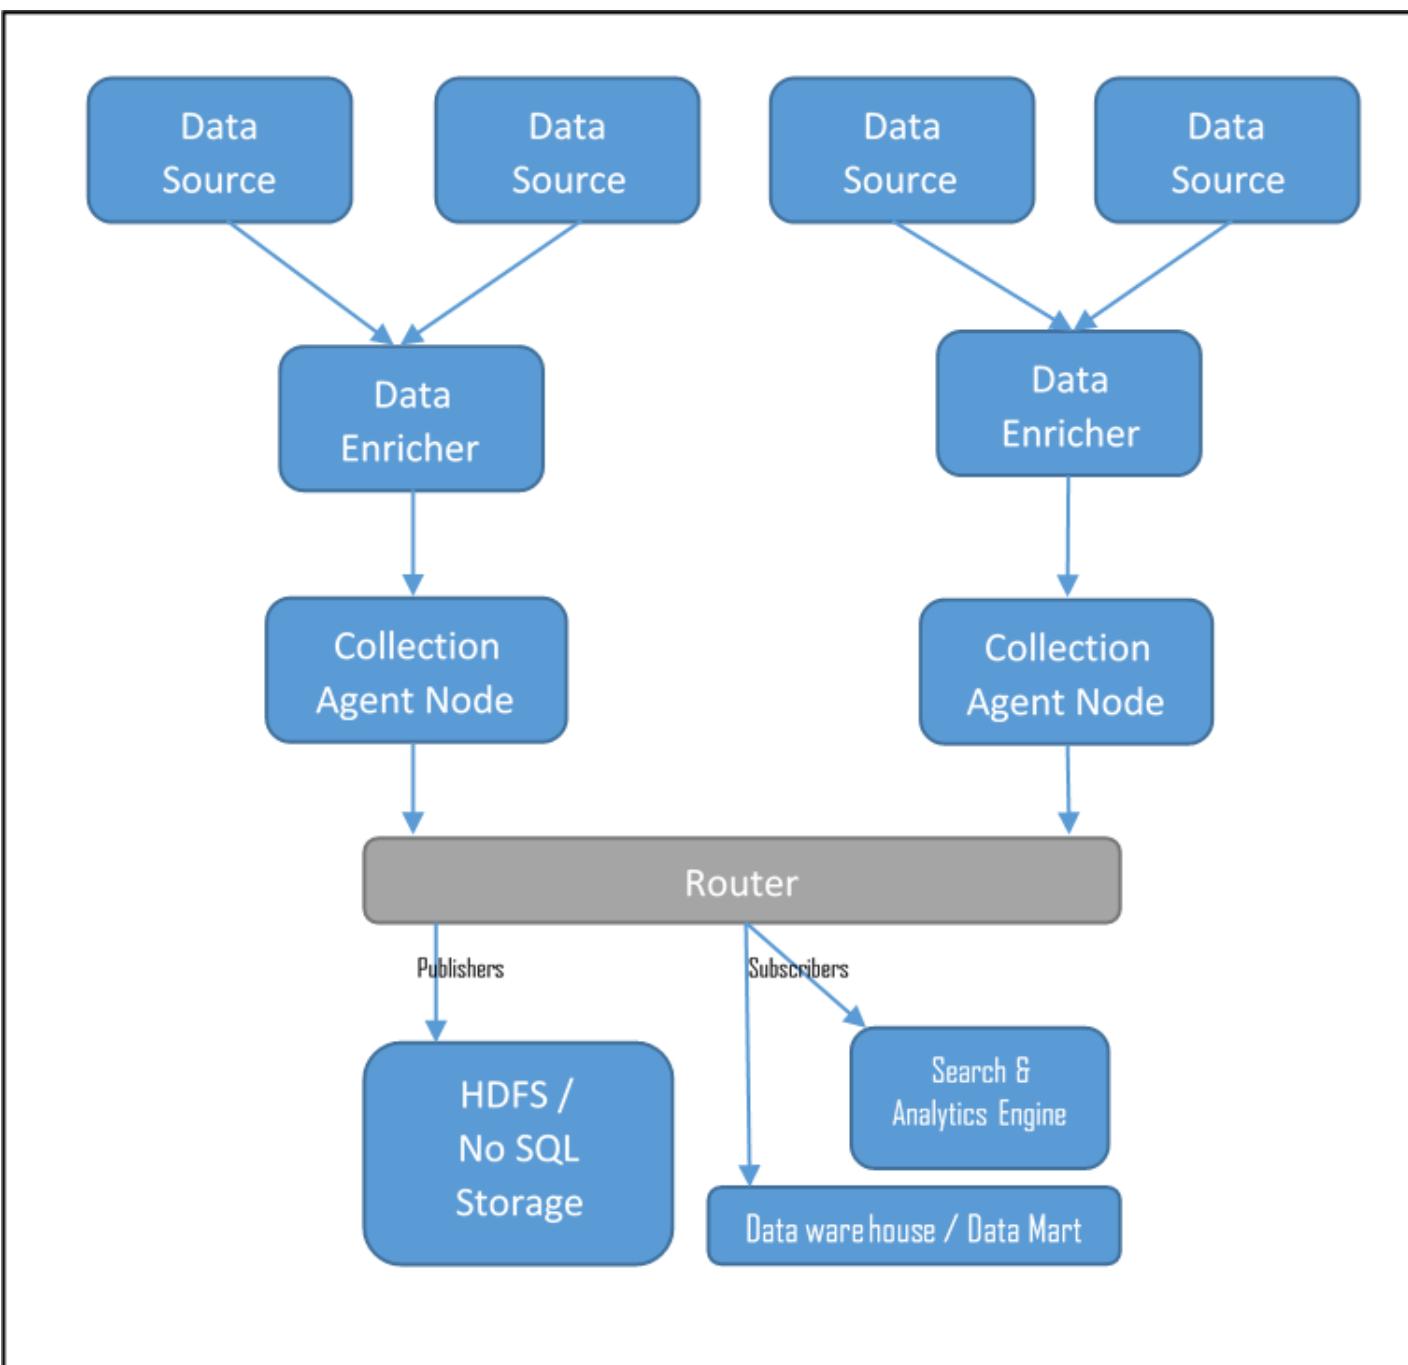
Ingestion and streaming patterns

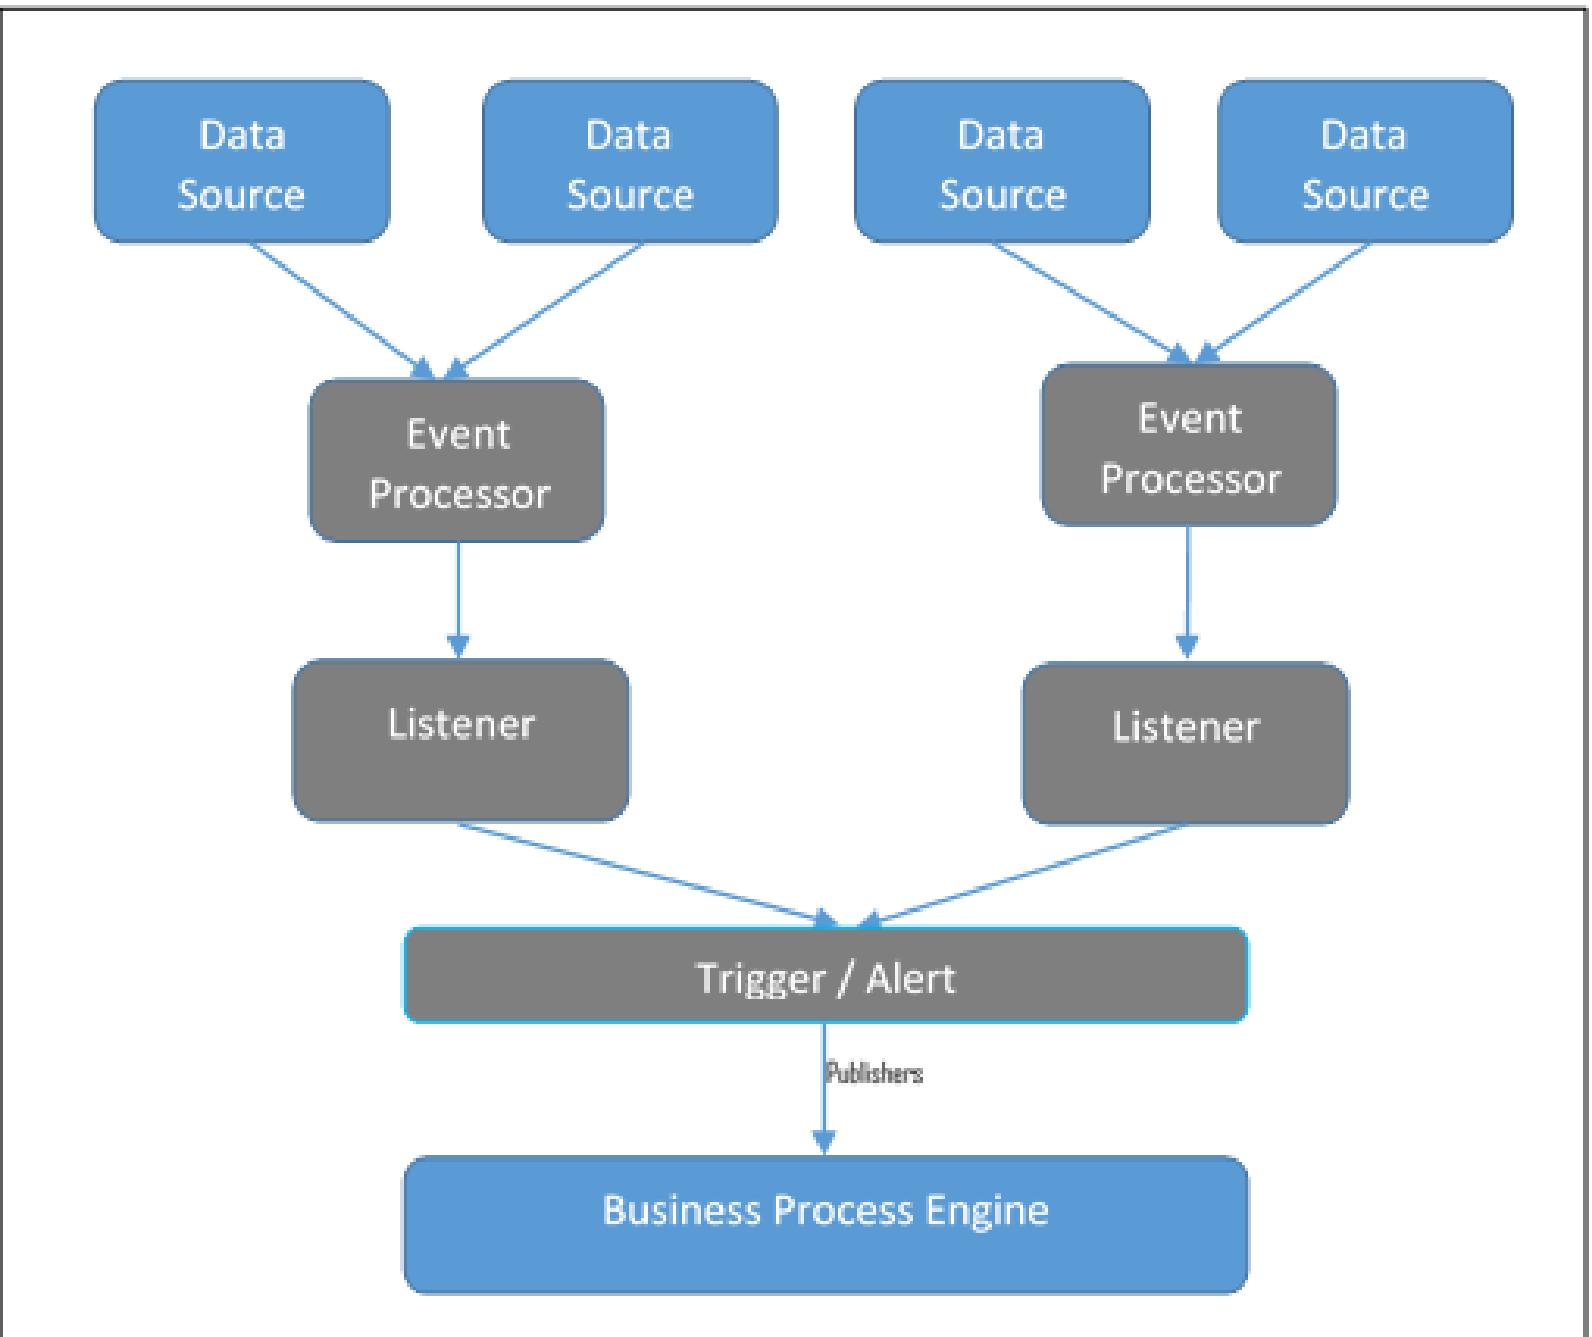
- Multisource extractor
- Multidestination
- Protocol converter
- **Just-in-time (JIT)**
transformation
- Real-time streaming pattern

Multisource extractor







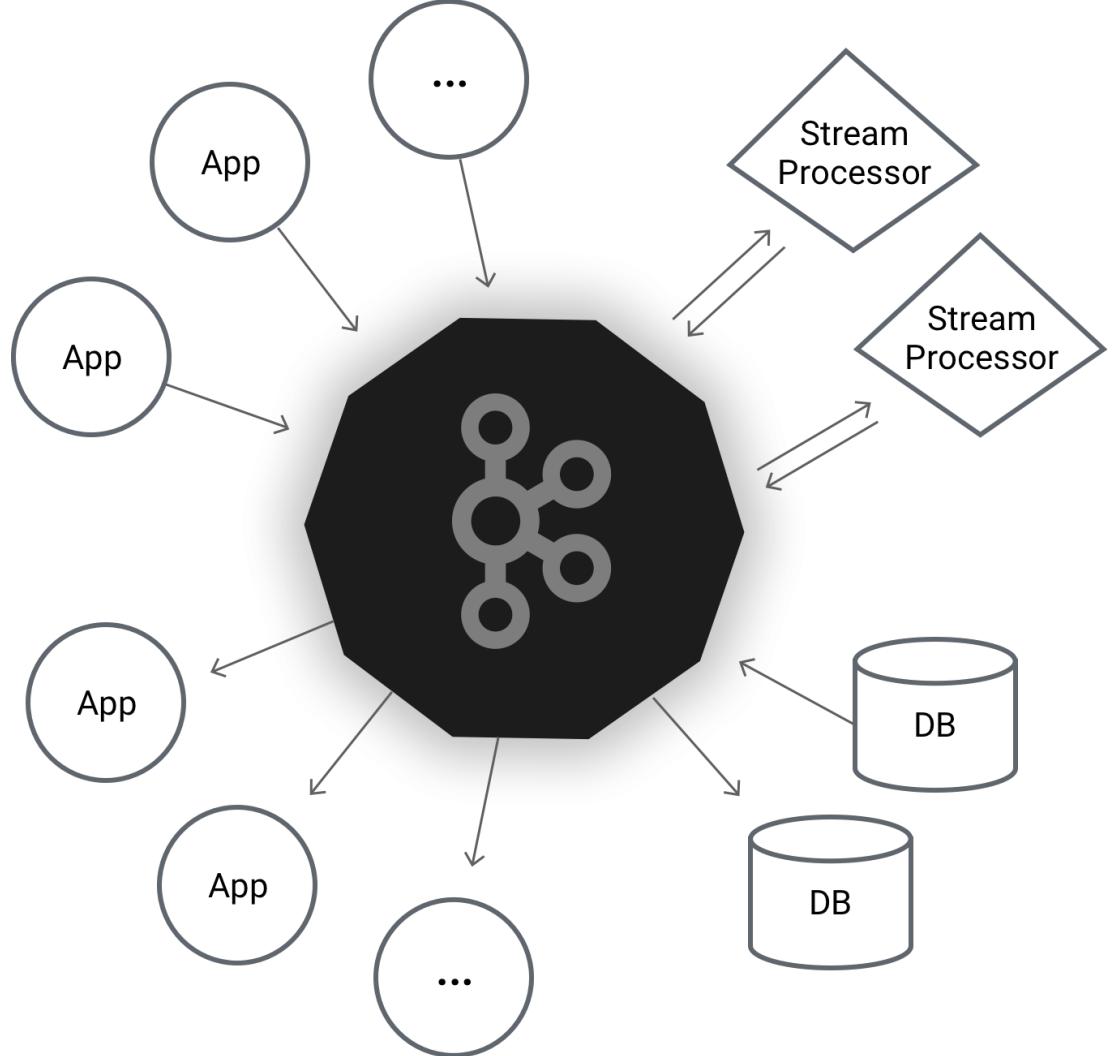


Real World Examples



A distributed streaming platform

<https://kafka.apache.org/>



Apache Kafka

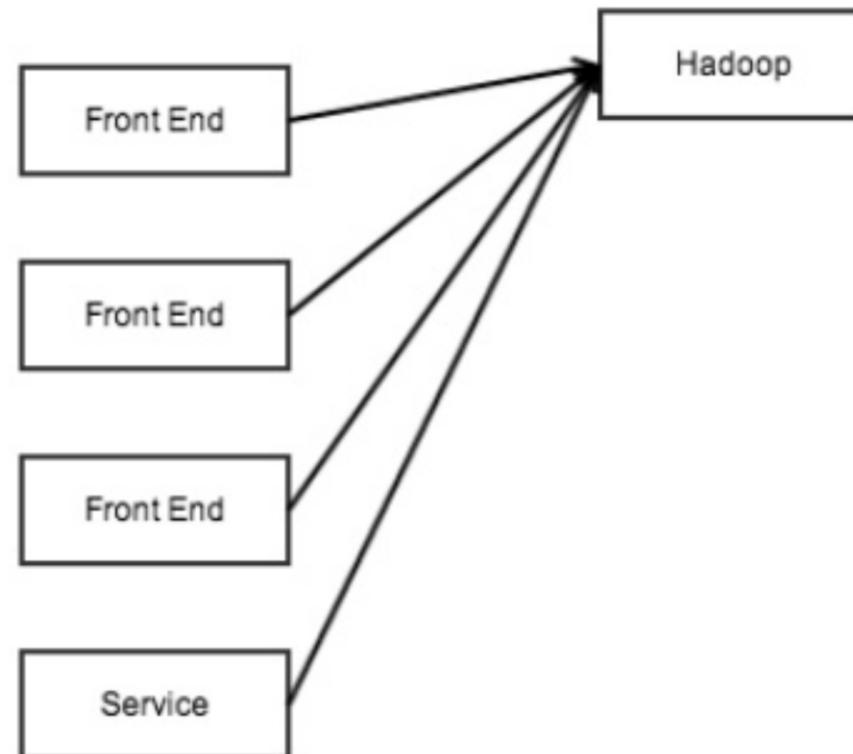
- What, Why, How of Apache Kafka
 - Producers, Brokers, Consumers, Topics and Partitions
- Get up and running - Quick Start
- Existing Integrations & Client Libraries
- Developing Producers
- Developing Consumers

Source: Joe Stein, <http://allthingsshadoop.com>

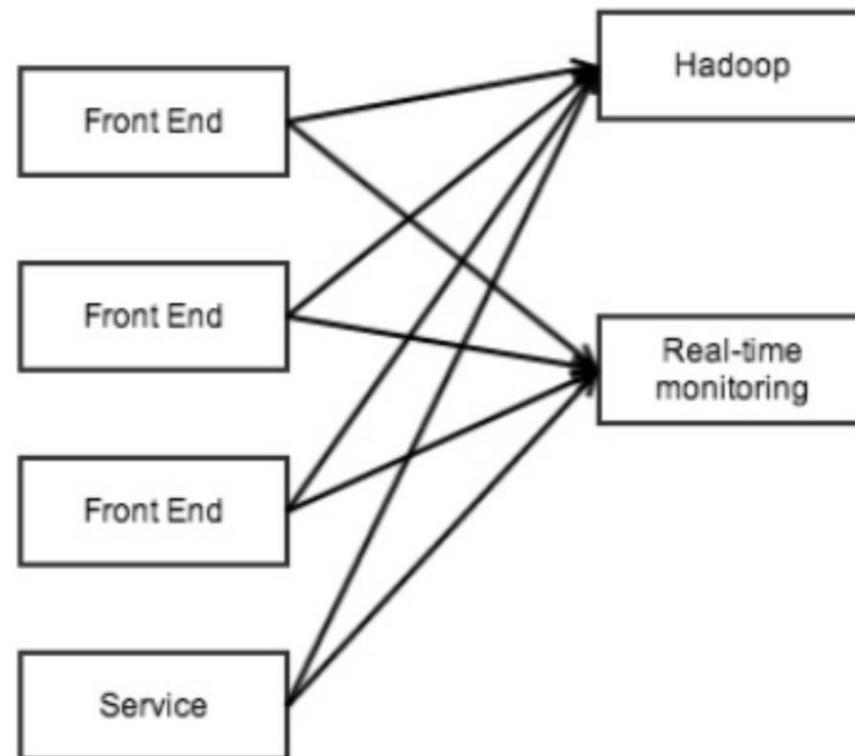
It often starts with just one data pipeline



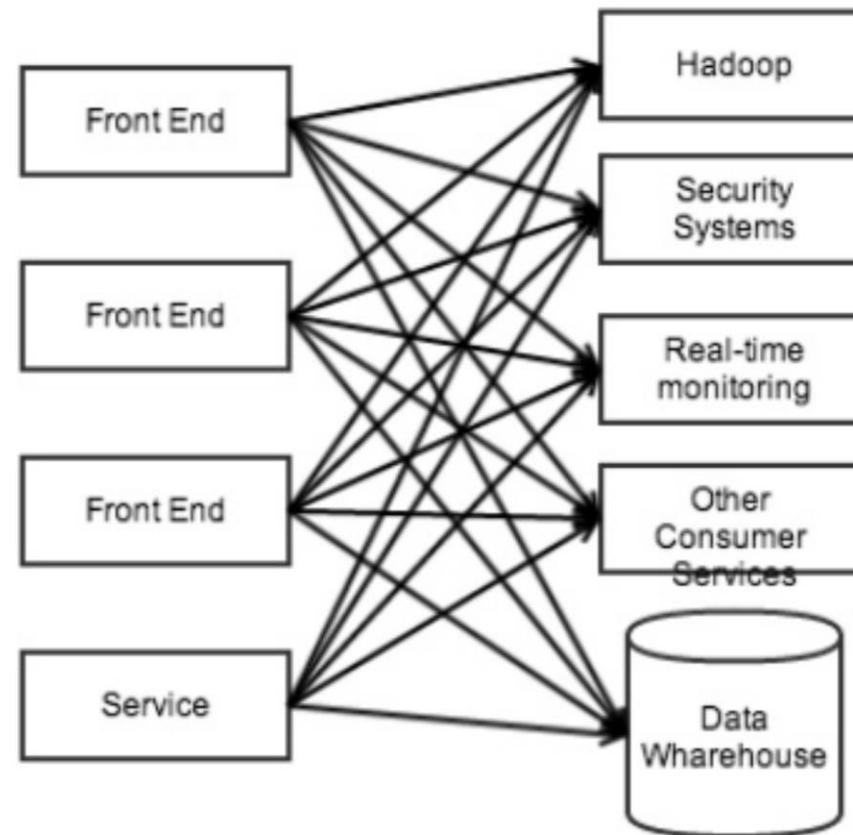
Reuse of data pipelines for new providers



Reuse of existing providers for new consumers

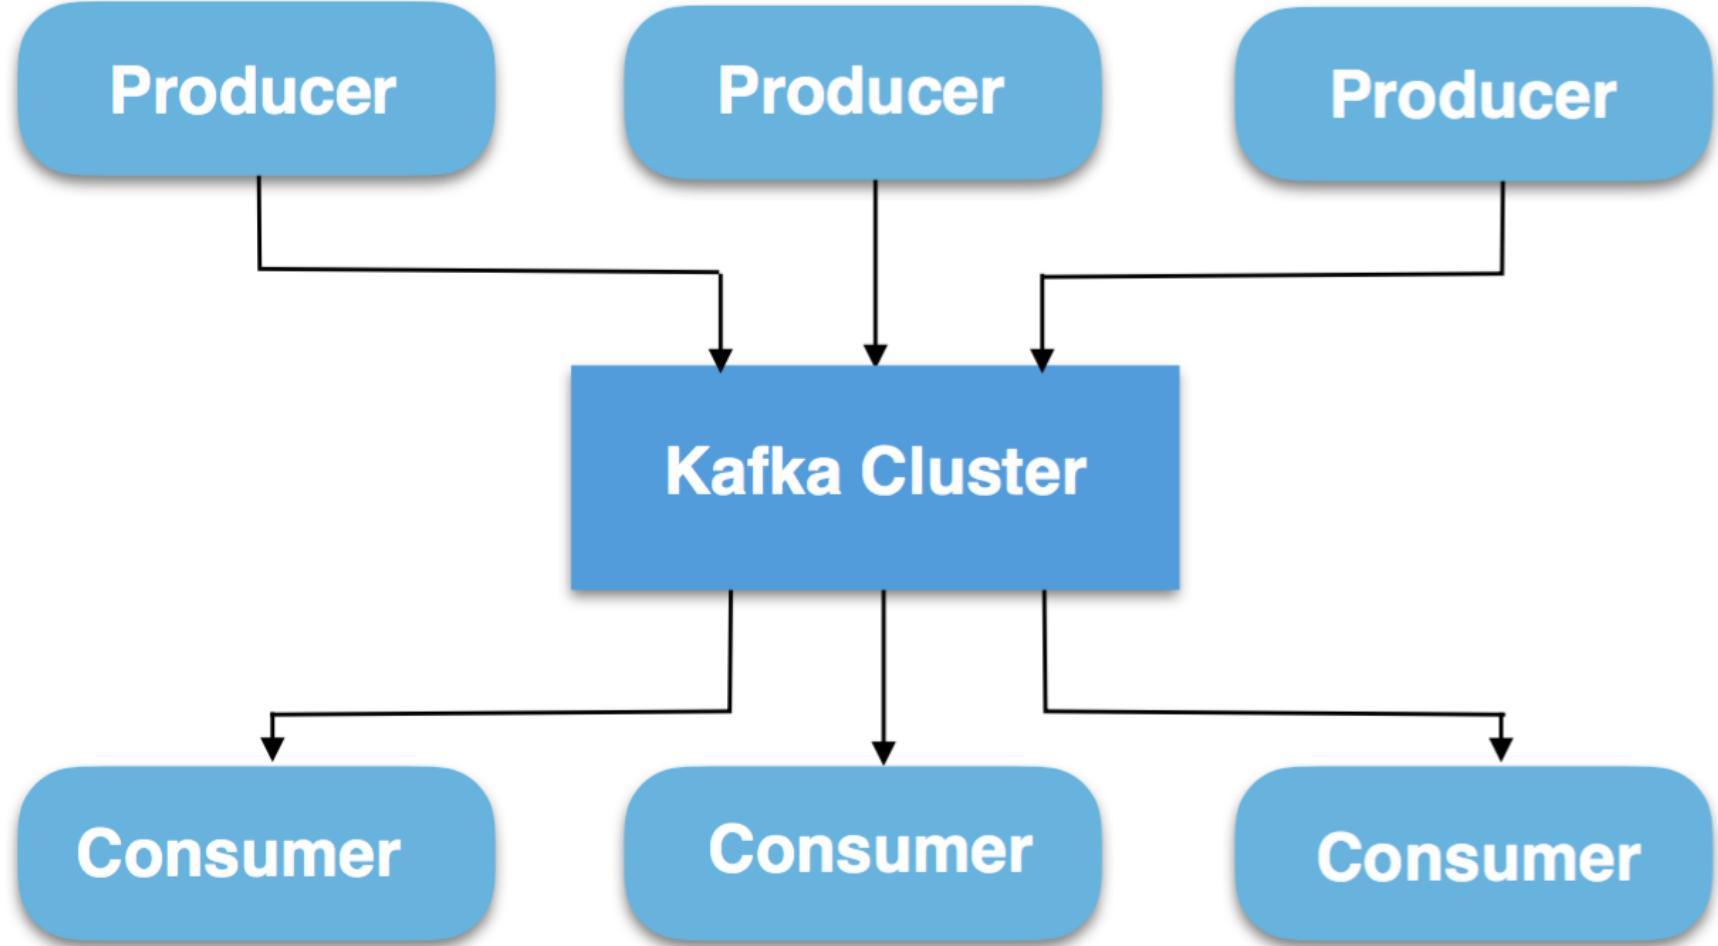


Eventually the solution becomes the problem



Apache Kafka

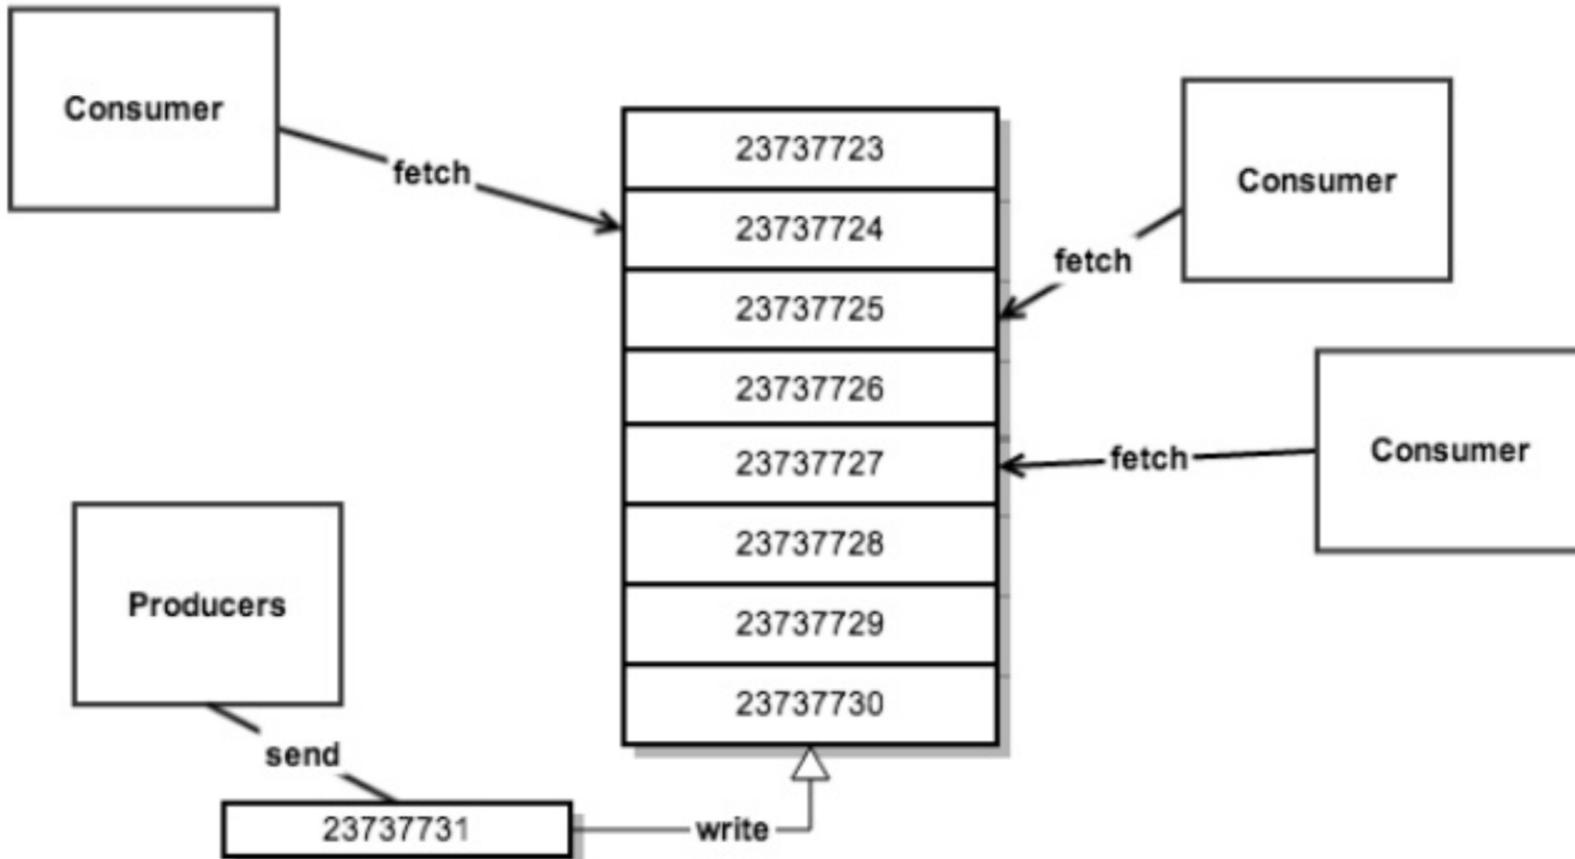
- Decouple the data pipeline
- Messaging System
 - Queues (Topics)
 - Writers (Producers)
 - Readers (Consumers)



How does Kafka do this?

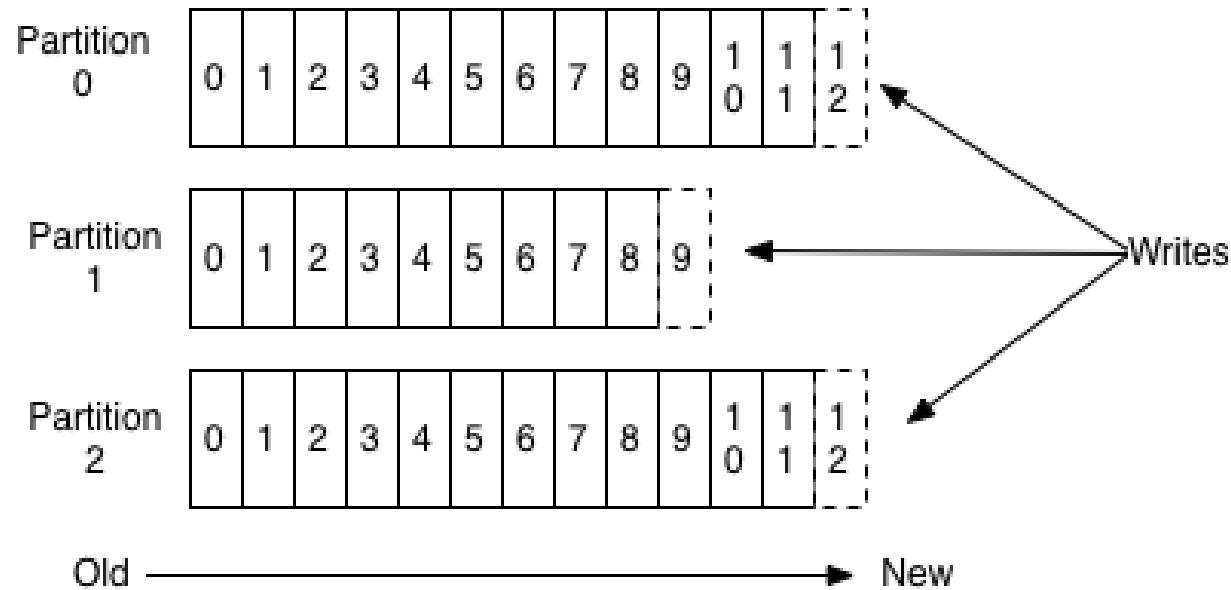
- Producers - ** push **
 - Batching
 - Compression
 - Sync (Ack), Async (auto batch)
 - Replication
 - Sequential writes, guaranteed ordering within each partition
- Consumers - ** pull **
 - No state held by broker
 - Consumers control reading from the stream
- Zero Copy for producers and consumers to and from the broker <http://kafka.apache.org/documentation.html#maximizingefficiency>
- Message stay on disk when consumed, deletes on TTL with compaction

A high-throughput distributed messaging system rethought as a distributed commit log.

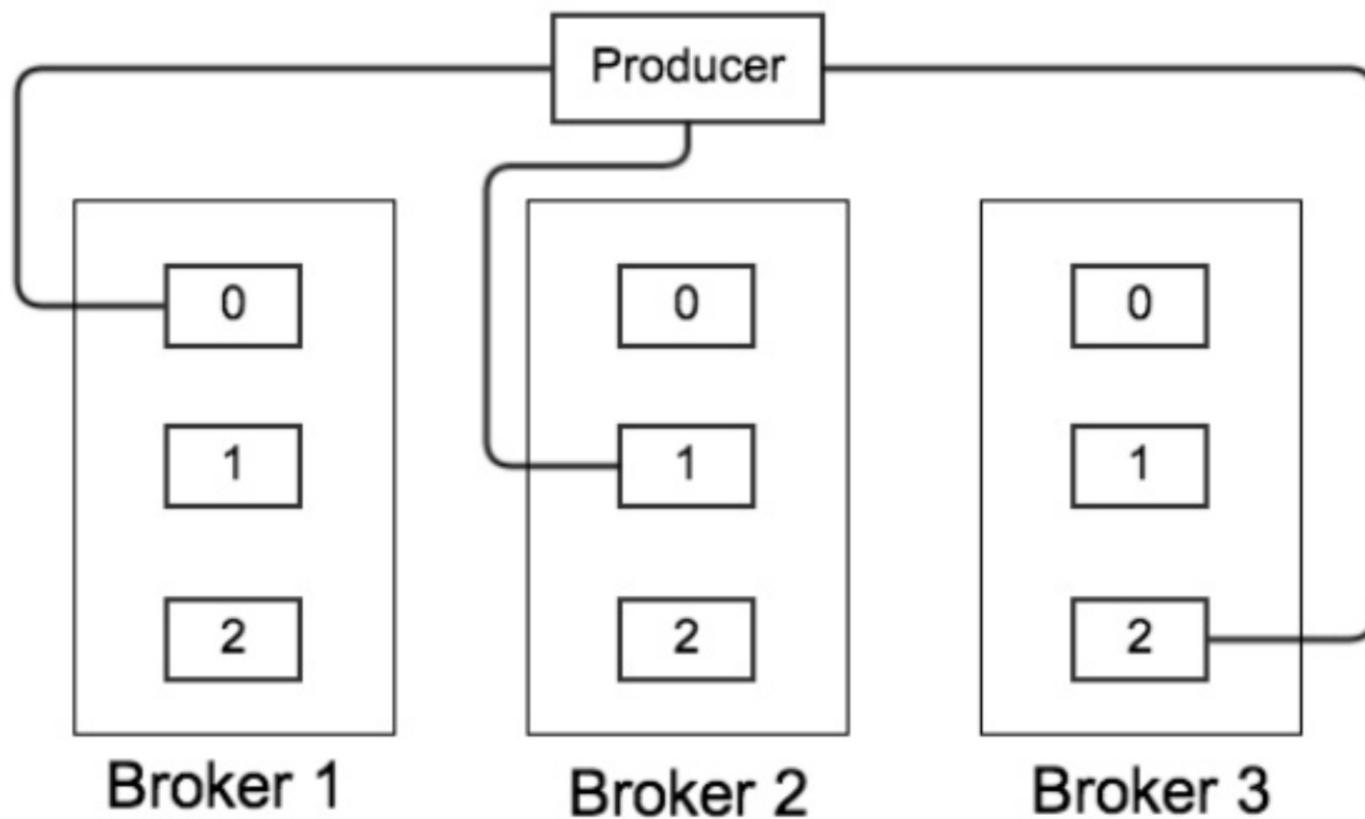


Apache Kafka

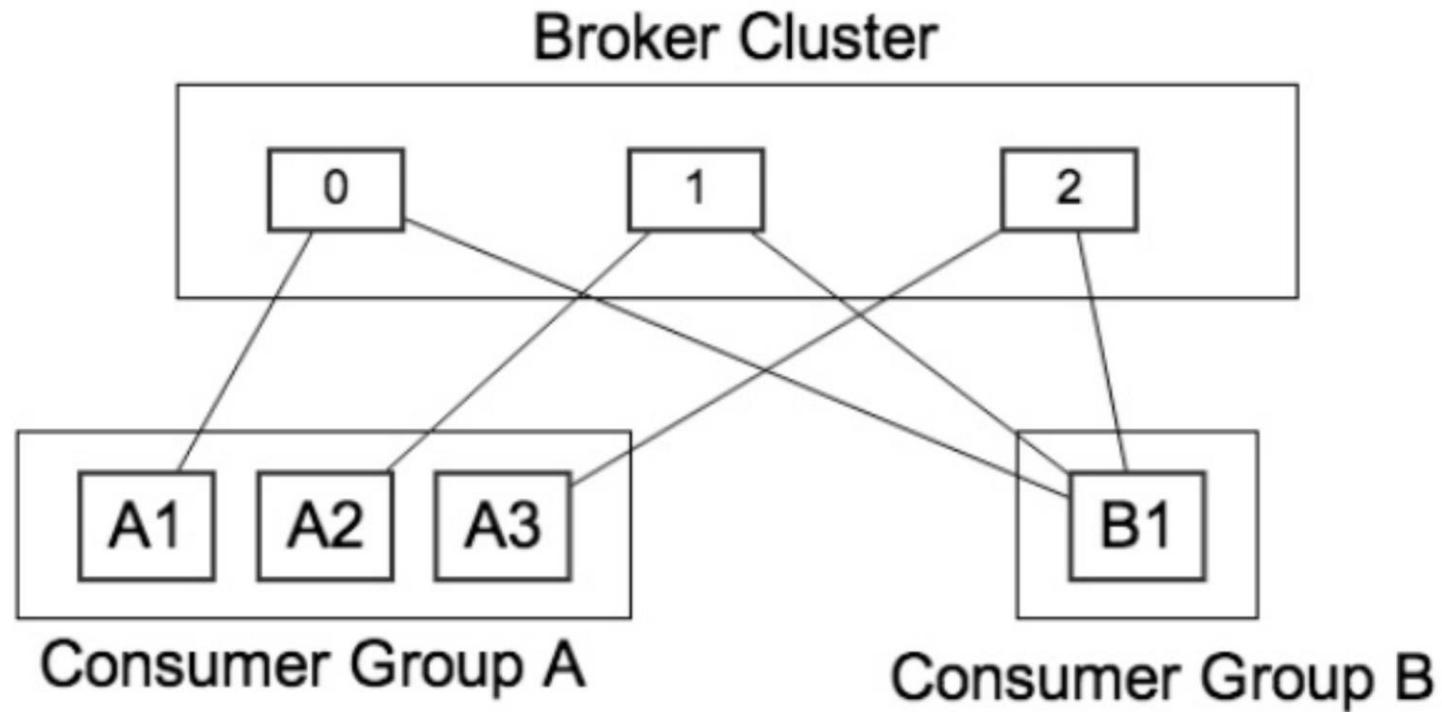
Anatomy of a Topic



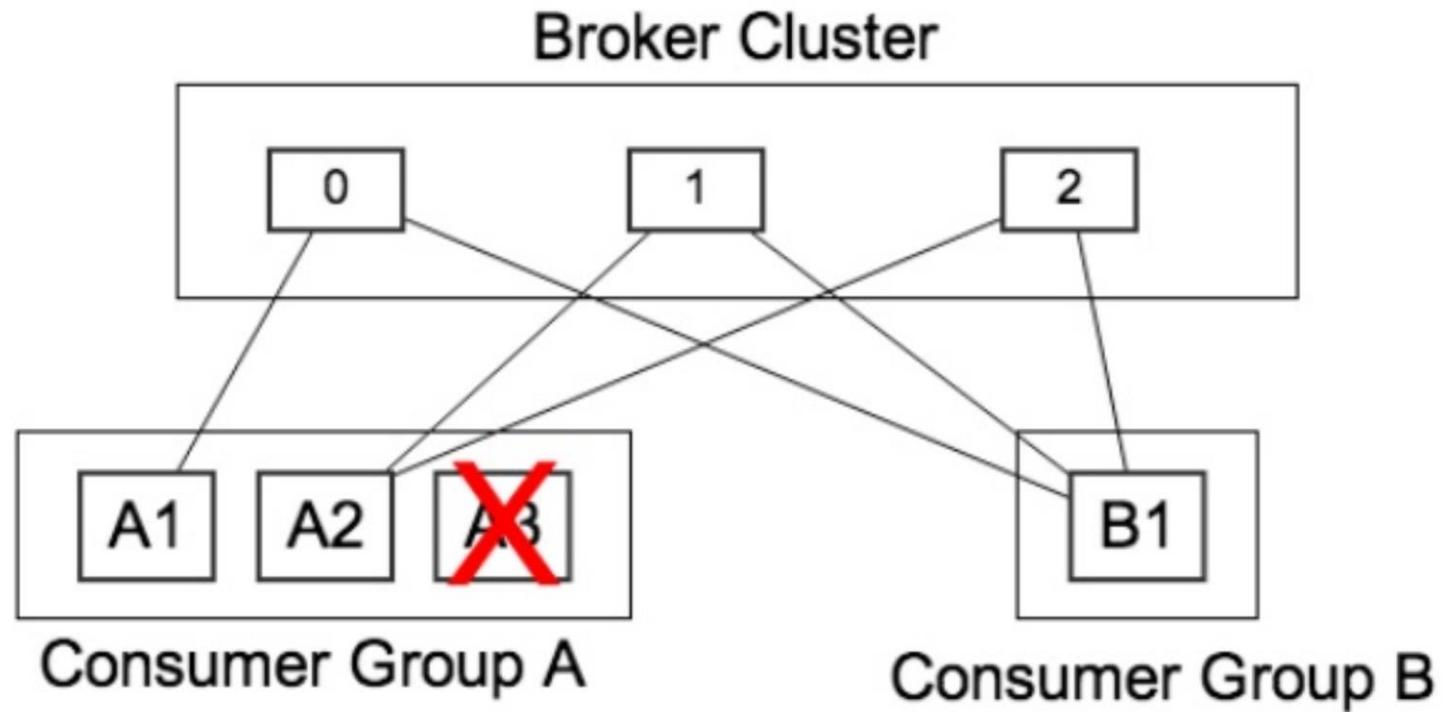
Brokers load balance producers by partition



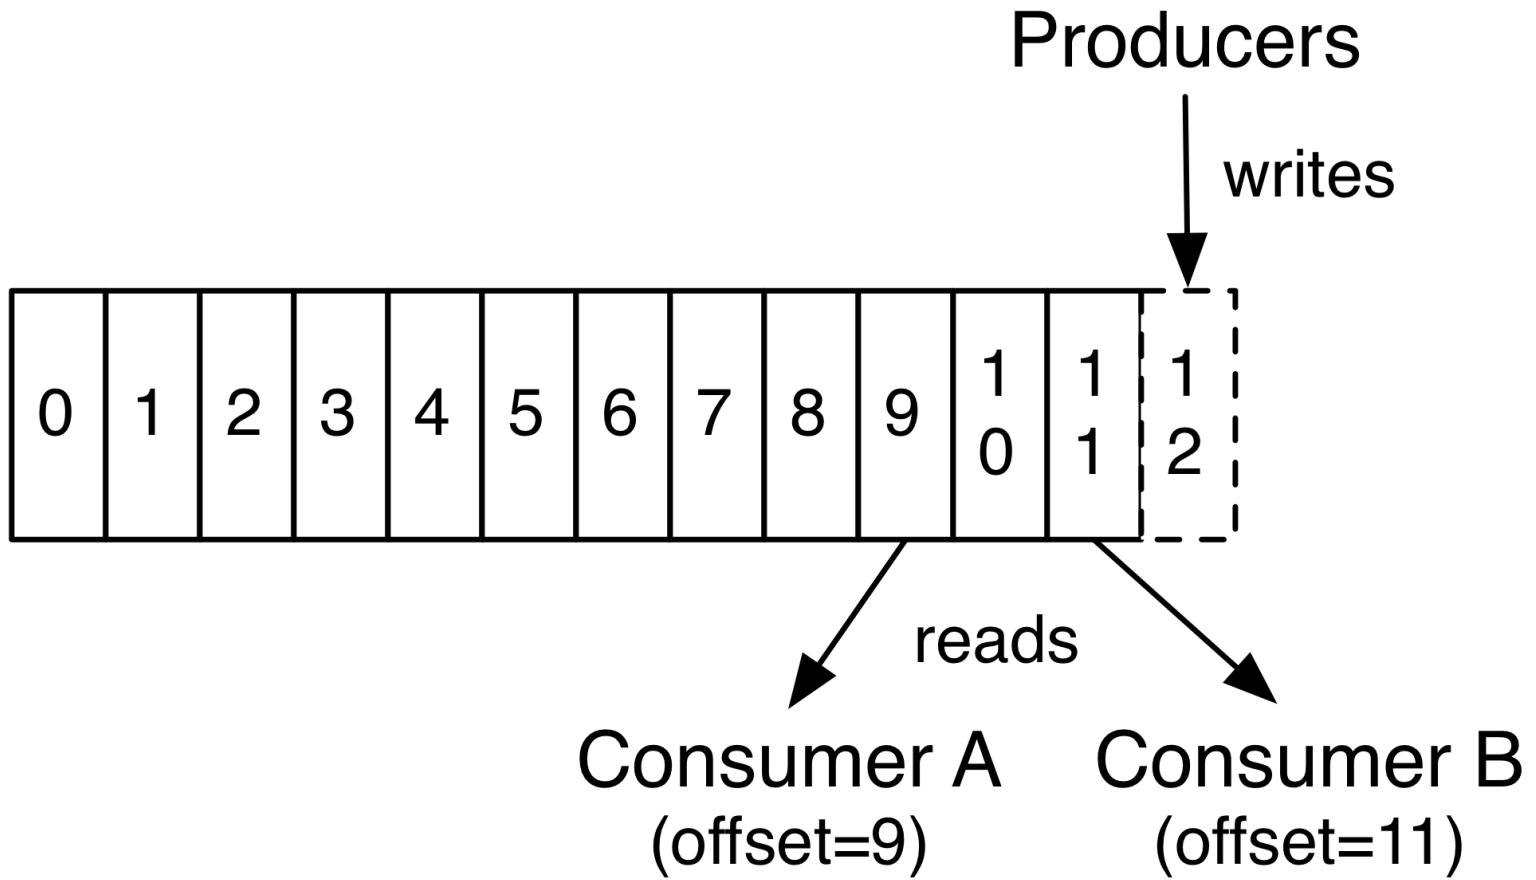
Consumer group provide isolation to topics and partitions



Consumer rebalance themselves for partitions



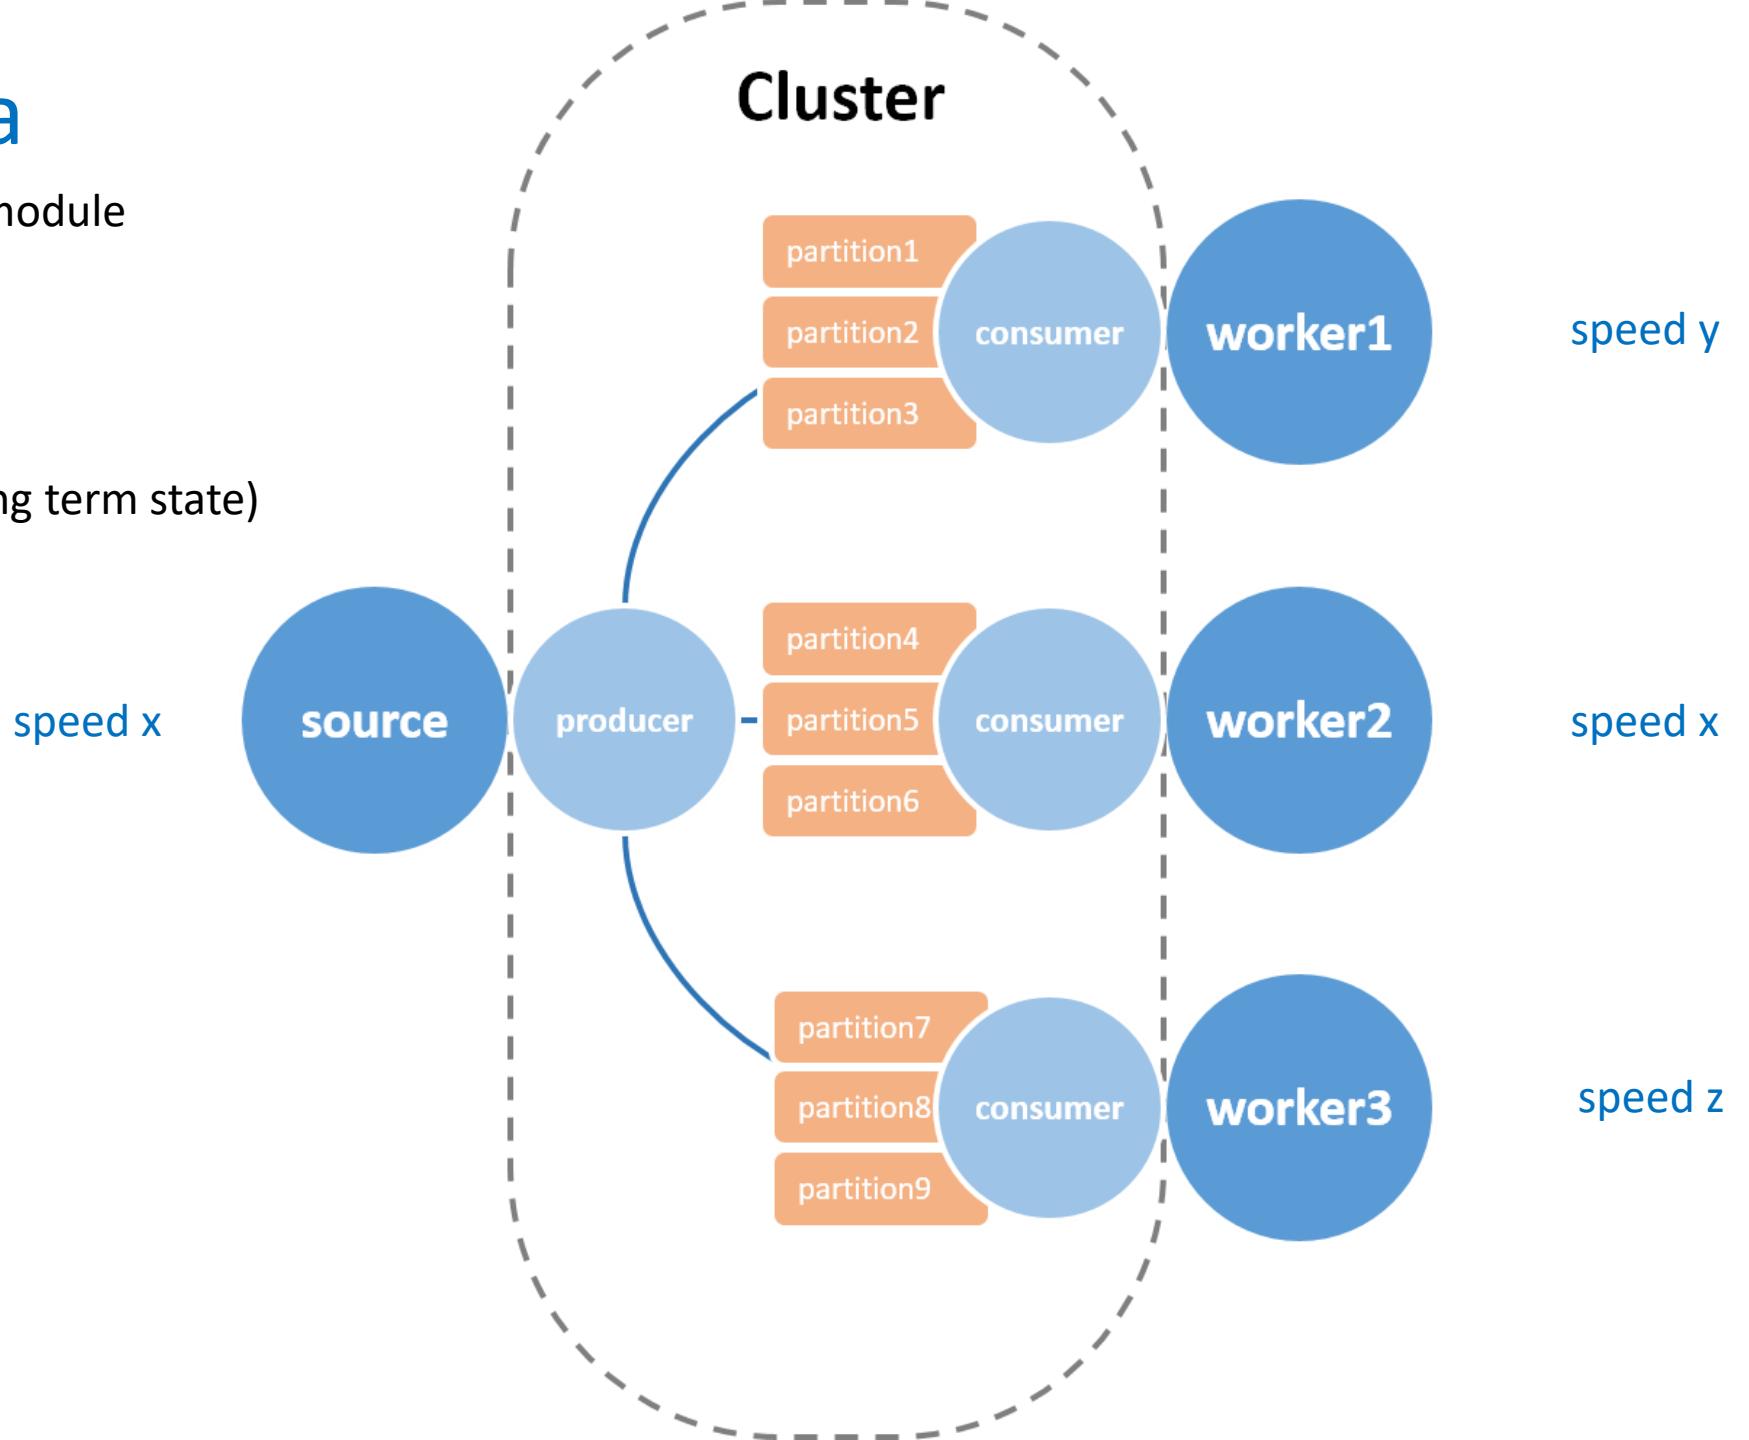
Apache Kafka



Apache Kafka

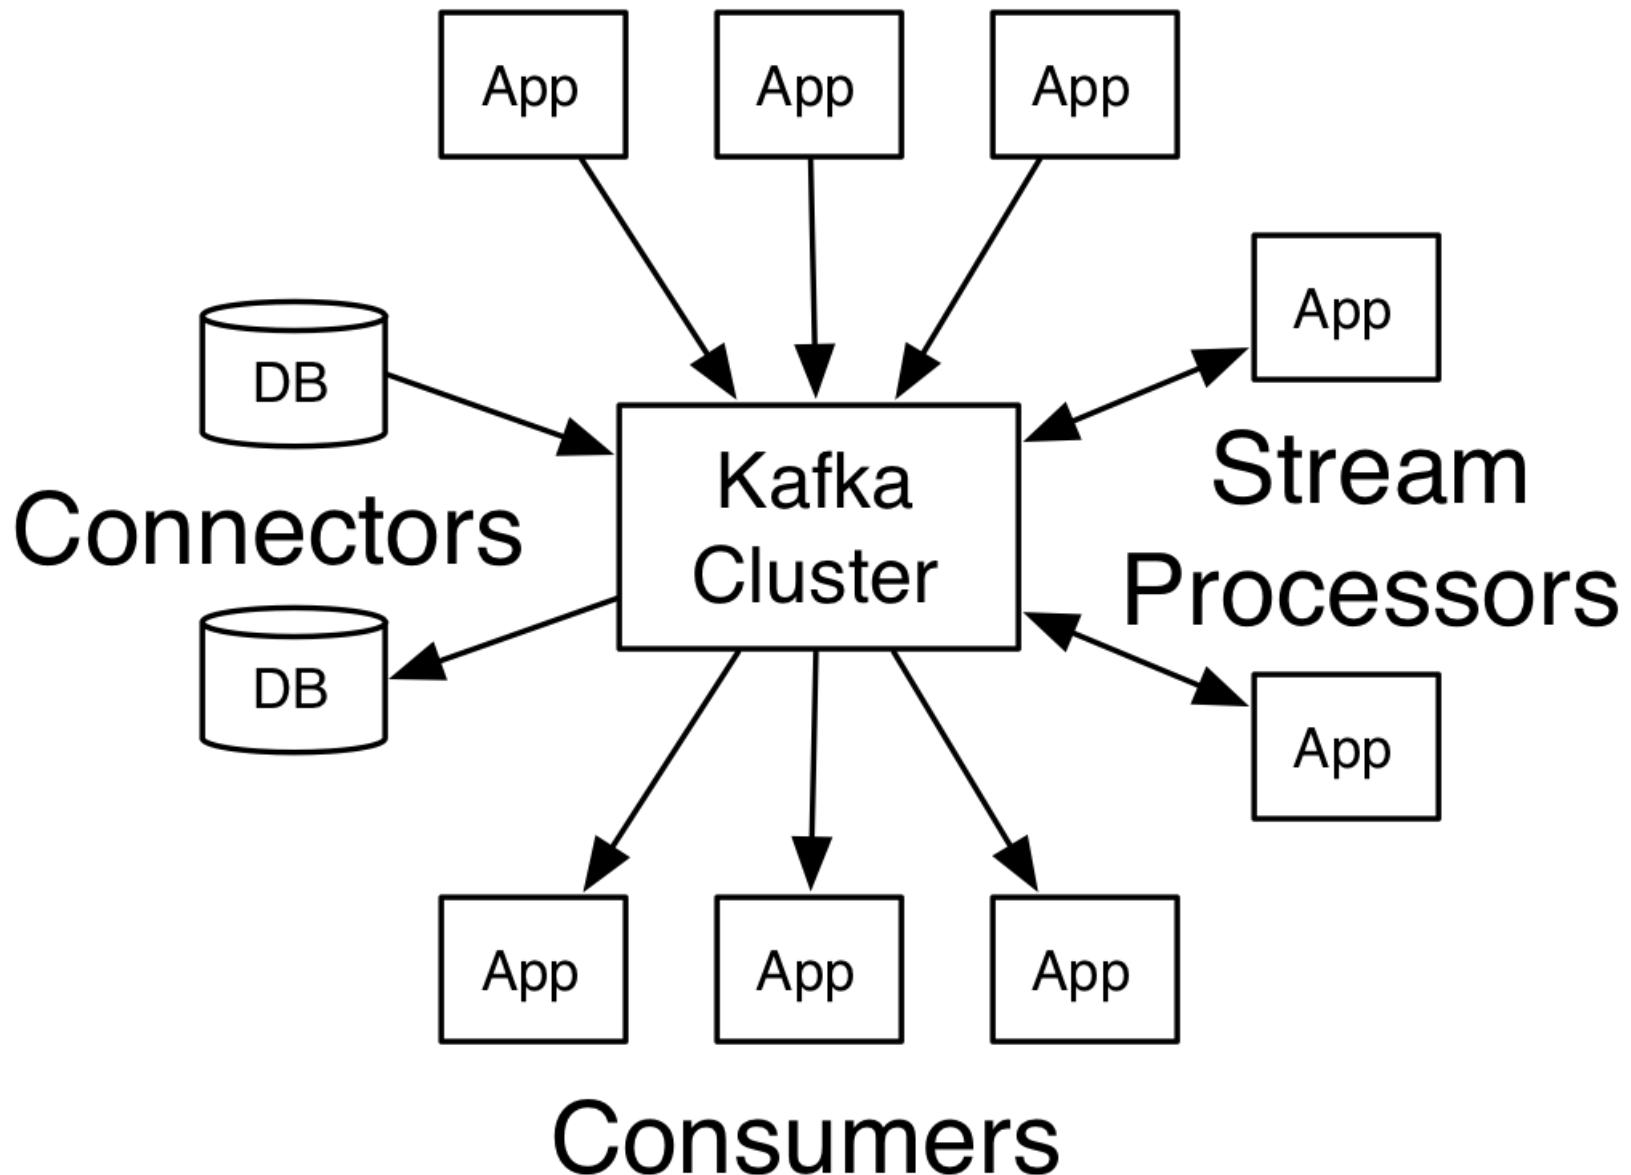
- Decoupling - remove module dependencies
- Asynchronous

$x \leq y, z$ in most cases (long term state)



Apache Kafka

Producers



Existing Integrations

<https://cwiki.apache.org/confluence/display/KAFKA/Ecosystem>

- log4j Appender
- Apache Storm
- Apache Camel
- Apache Samza
- Apache Hadoop
- Apache Flume
- Camus
- AWS S3
- Rieman
- Sematext
- Dropwizard
- LogStash
- Fluent

Client Libraries

Community Clients <https://cwiki.apache.org/confluence/display/KAFKA/Clients>

- Python - Pure Python implementation with full protocol support. Consumer and Producer implementations included, GZIP and Snappy compression supported.
- C - High performance C library with full protocol support
- C++ - Native C++ library with protocol support for Metadata, Produce, Fetch, and Offset.
- Go (aka golang) Pure Go implementation with full protocol support. Consumer and Producer implementations included, GZIP and Snappy compression supported.
- Ruby - Pure Ruby, Consumer and Producer implementations included, GZIP and Snappy compression supported. Ruby 1.9.3 and up (CI runs MRI 2).
- Clojure - Clojure DSL for the Kafka API
- JavaScript (NodeJS) - NodeJS client in a pure JavaScript implementation

Wire Protocol Developers Guide

<https://cwiki.apache.org/confluence/display/KAFKA/A+Guide+To+The+Kafka+Protocol>

Developing Producers

<https://github.com/stealthly/scala-kafka/blob/master/src/test/scala/KafkaSpec.scala>

```
val producer = new KafkaProducer("test-topic","192.168.86.10:9092")
producer.send("hello distributed commit log")
```

Producers

<https://github.com/stealthly/scala-kafka/blob/master/src/main/scala/KafkaProducer.scala>

```
case class KafkaProducer(
```

```
    topic: String,
```

```
    brokerList: String,
```

```
    /** brokerList - This is for bootstrapping and the producer will only use it for  
     * getting metadata (topics, partitions and replicas). The socket connections for  
     * sending the actual data will be established based on the broker information  
     * returned in the metadata. The format is host1:port1,host2:port2, and the list can  
     * be a subset of brokers or a VIP pointing to a subset of brokers.
```

```
    */
```

Producer

```
clientId: String = UUID.randomUUID().toString,
```

```
/** clientId - The client id is a user-specified string sent in each request to help  
trace calls. It should logically identify the application making the request. */
```

```
synchronously: Boolean = true,
```

```
/** synchronously - This parameter specifies whether the messages are sent  
asynchronously in a background thread. Valid values are false for  
asynchronous send and true for synchronous send. By setting the producer to  
async we allow batching together of requests (which is great for throughput) but  
open the possibility of a failure of the client machine dropping unsent data. */
```

Producer

`compress: Boolean = true,`

`/** compress -This parameter allows you to specify the compression codec for all data generated by this producer. When set to true gzip is used. To override and use snappy you need to implement that as the default codec for compression using SnappyCompressionCodec.codec instead of DefaultCompressionCodec.codec below. */`

`batchSize: Integer = 200,`

`/** batchSize -The number of messages to send in one batch when using async mode. The producer will wait until either this number of messages are ready to send or queue.buffer.max.ms is reached.*/`

Producer

```
messageSendMaxRetries: Integer = 3,
```

```
/** messageSendMaxRetries - This property will cause the producer to  
automatically retry a failed send request. This property specifies the number of  
retries when such failures occur. Note that setting a non-zero value here can  
lead to duplicates in the case of network errors that cause a message to be  
sent but the acknowledgement to be lost.*/
```

Producer

```
requestRequiredAcks: Integer = -1
```

```
/** requestRequiredAcks
```

0) which means that the producer never waits for an acknowledgement from the broker (the same behavior as 0.7). This option provides the lowest latency but the weakest durability guarantees (some data will be lost when a server fails).

1) which means that the producer gets an acknowledgement after the leader replica has received the data. This option provides better durability as the client waits until the server acknowledges the request as successful (only messages that were written to the now-dead leader but not yet replicated will be lost).

-1) which means that the producer gets an acknowledgement after all in-sync replicas have received the data. This option provides the best durability, we guarantee that no messages will be lost as long as at least one in sync replica remains.*/

High Level Consumer

<https://github.com/stealthly/scala-kafka/blob/master/src/main/scala/KafkaConsumer.scala>

```
class KafkaConsumer(
```

```
topic: String,
```

```
/** topic - The high-level API hides the details of brokers from the consumer  
and allows consuming off the cluster of machines without concern for the  
underlying topology. It also maintains the state of what has been consumed.  
The high-level API also provides the ability to subscribe to topics that match a  
filter expression (i.e., either a whitelist or a blacklist regular expression).*/
```

High Level Consumer

```
groupId: String,  
/** groupId - A string that uniquely identifies the group of consumer processes  
to which this consumer belongs. By setting the same group id multiple  
processes indicate that they are all part of the same consumer group.*/  
zookeeperConnect: String,  
/** zookeeperConnect - Specifies the zookeeper connection string in the form  
hostname:port where host and port are the host and port of a zookeeper server.  
To allow connecting through other zookeeper nodes when that zookeeper  
machine is down you can also specify multiple hosts in the form hostname1:  
port1,hostname2:port2,hostname3:port3. The server may also have a  
zookeeper chroot path as part of it's zookeeper connection string which puts its  
data under some path in the global zookeeper namespace. */
```

Simple Consumer

<https://cwiki.apache.org/confluence/display/KAFKA/0.8.0+SimpleConsumer+Example>

<https://github.com/apache/kafka/blob/0.8/core/src/main/scala/kafka/tools/SimpleConsumerShell.scala>

```
val fetchRequest = fetchRequestBuilder  
    .addFetch(topic, partitionId, offset, fetchSize)  
    .build()
```

Apache Kafka

Tutorial – Docker based

<https://docs.confluent.io/current/quickstart/ce-docker-quickstart.html>