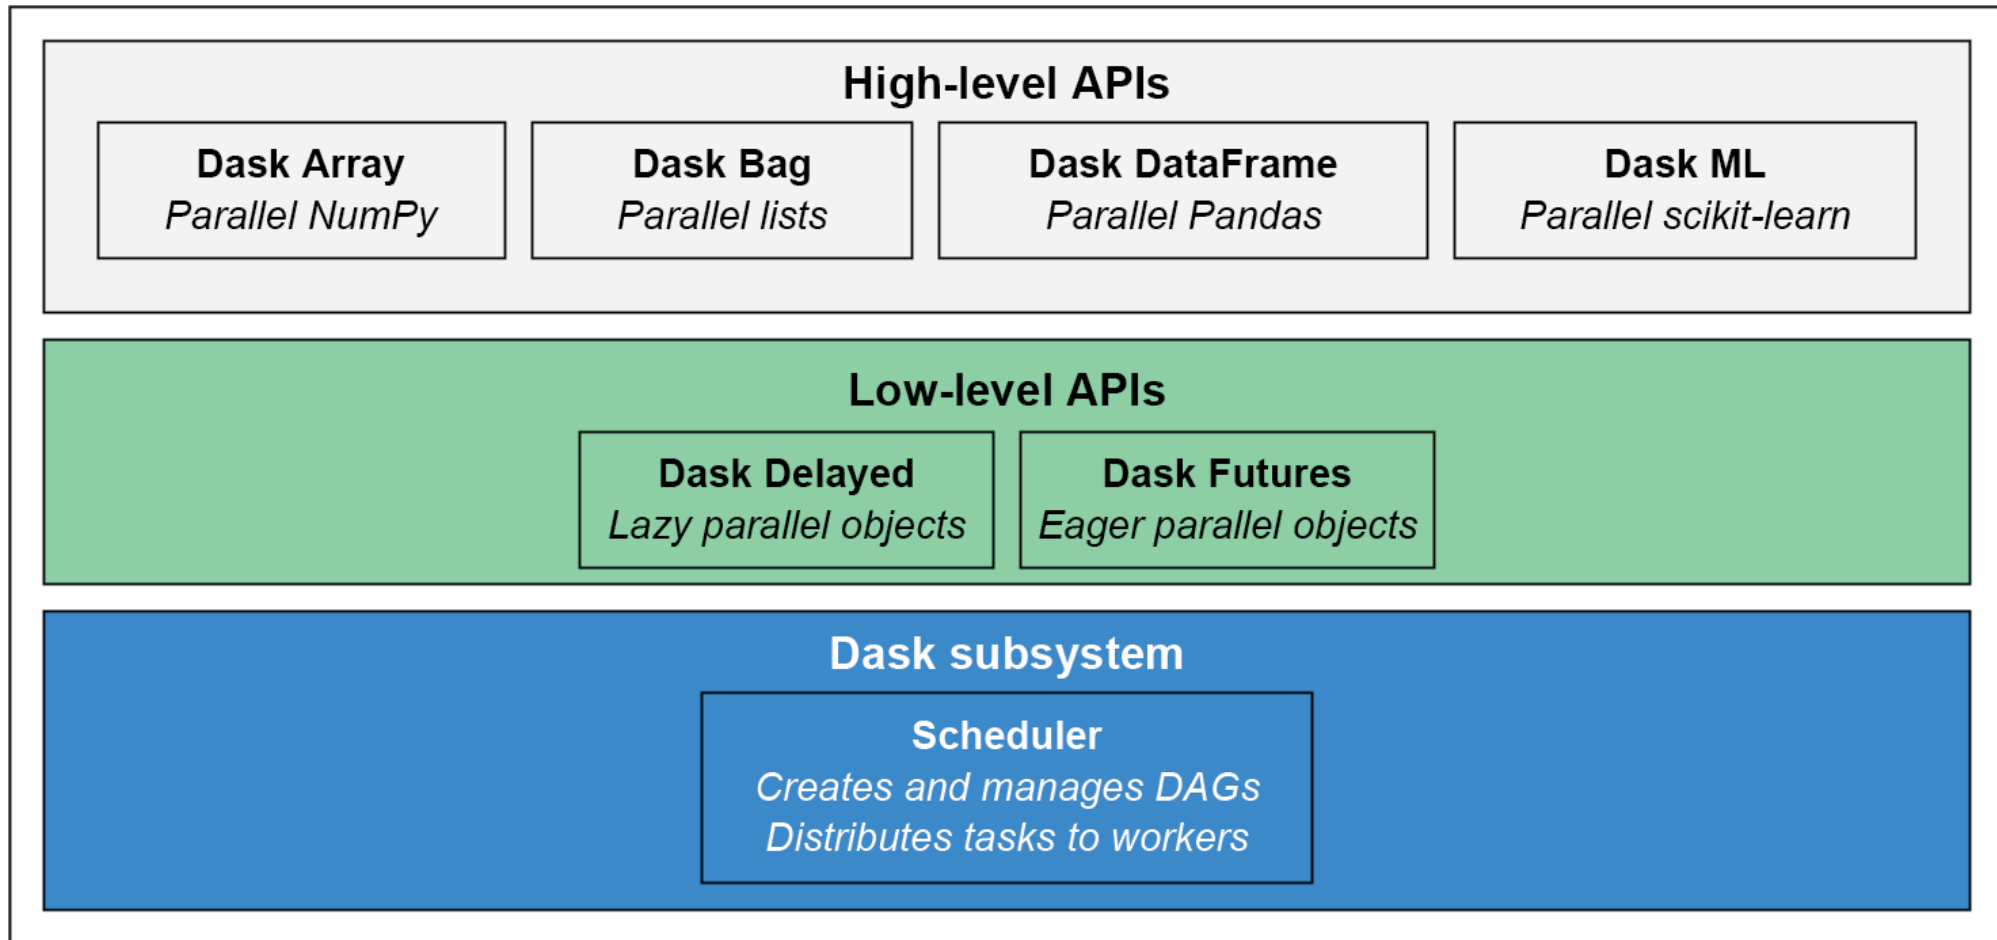


Python & Dask

Parallel & Distributed Computing with Python, Pandas and NumPy

Dask: <https://docs.dask.org/en/stable/>



Collections

(create task graphs)



Task Graph



Schedulers

(execute task graphs)


Dask Array

Dask DataFrame

Dask Bag


Dask Delayed

Futures



Single-machine
(threads, processes,
synchronous)

Distributed



Dask DataFrame structure:

	Summons number	Plate ID	Registration state	Plate type	Issue date	Violation code	Vehicle body type	Vehicle make	Issuing agency	Street code1	Street code2	Street code3	Vehicle expiration date	Violation location	Violation precinct
npartitions=33	int64	object	object	object	object	int64	object	object	object	int64	int64	int64	int64	float64	int64

...

Dask Name: from-delayed, 99 tasks

Dask DataFrame structure

Column names

Summons
number

Plate
ID

Registration
state

Plate
type

Issue
date

Violation
code

Number of
partitions

npartitions=33

Column
datatypes

int64

object

object

object

object

int64

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

Dask name: from-delayed, 99 tasks

Internal name of the Number of nodes in the
underlying DAG

underlying DAG



Axis 1 columns

	Person ID	Last name	First name	Date of birth
0	1	Smith	John	10/6/82
1	2	Williams	Bill	7/4/90
2	3	Williams	Jane	5/6/89

Axis 0 rows

Default axis for
DataFrame operations

Index

- Provides an identifier for each row
- Defaults to sequential integers
- Used for grouping and joining operations

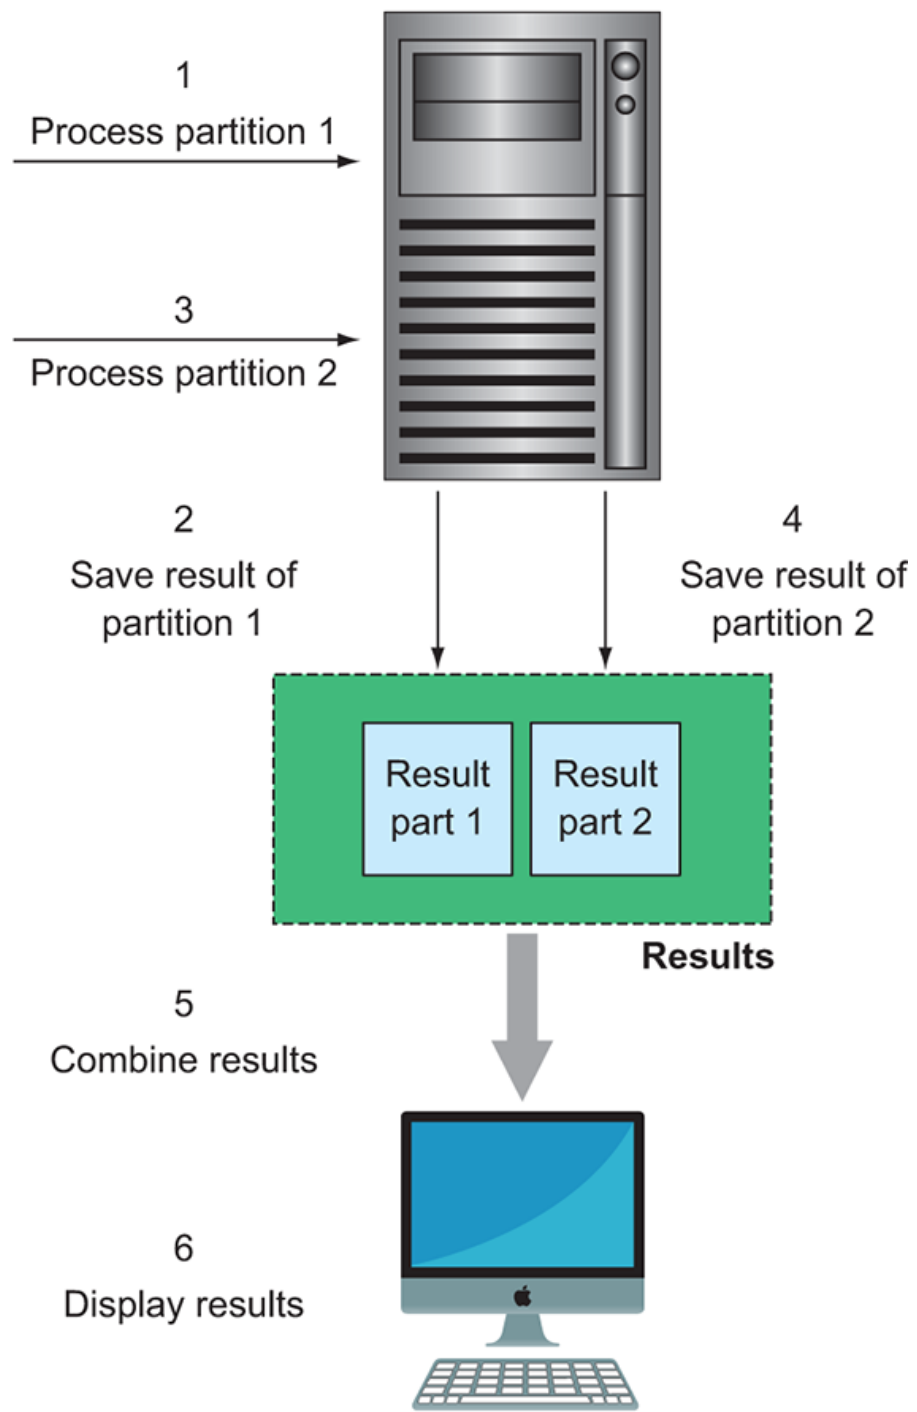
Dask DataFrame

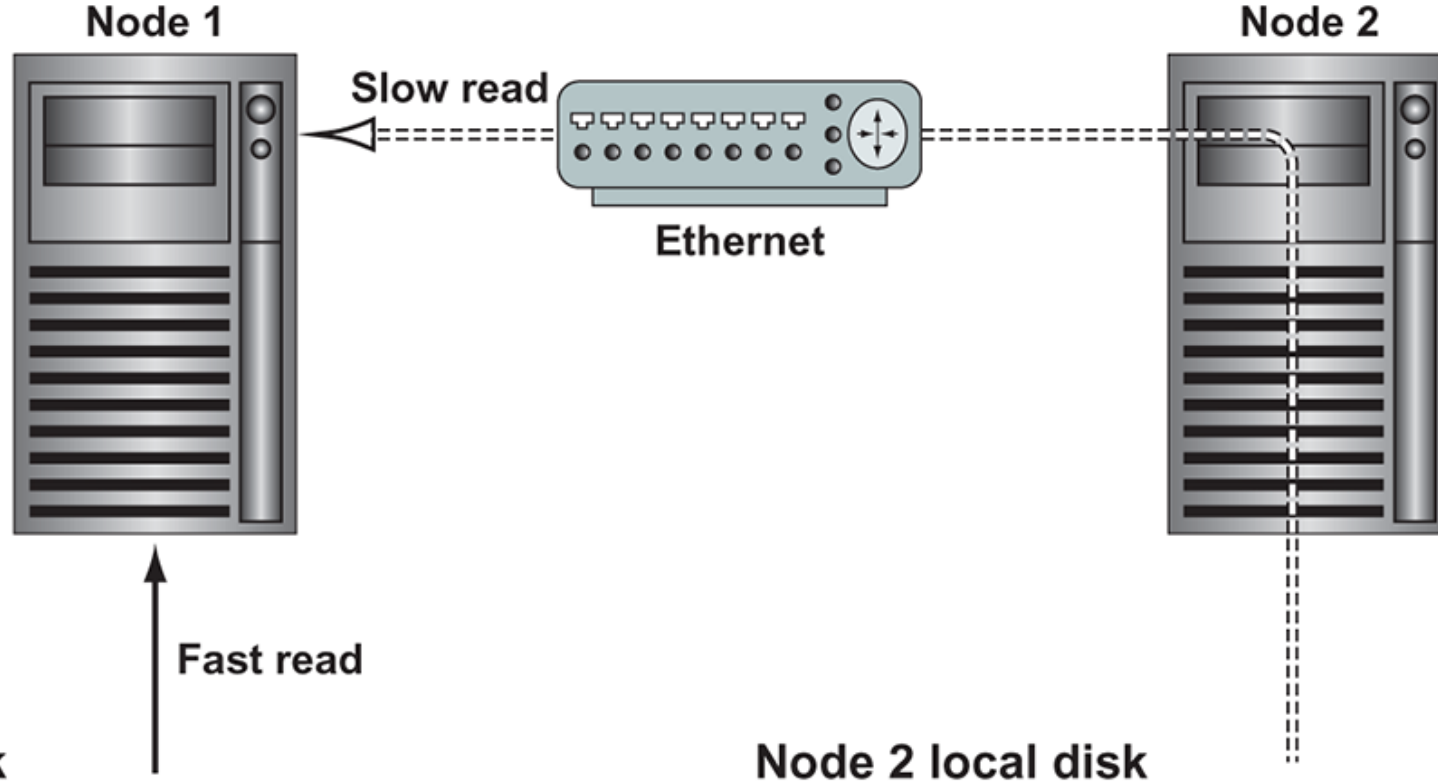
Partition 1

	Person ID	Last name	First name	Date of birth
0	1	Smith	John	10/6/82
1	2	Williams	Bill	7/4/90
2	3	Williams	Jane	5/6/89
3	4	Jackson	Cathy	1/24/74
4	5	Johnson	Stuart	6/5/95

Partition 2

	Person ID	Last name	First name	Date of birth
5	6	Smith	James	4/16/84
6	7	Anderson	Felicity	9/15/76
7	8	Christiansen	Liam	10/2/92
8	9	Carter	Nancy	2/5/86
9	10	Davidson	Christina	8/11/93





Node 1 local disk

Partition 1				
	Person ID	Last name	First name	Date of birth
0	1	Smith	John	10/6/82
1	2	Williams	Bill	7/4/90
2	3	Williams	Jane	5/6/89
3	4	Jackson	Cathy	1/24/74
4	5	Johnson	Stuart	6/5/95

Node 2 local disk

Partition 1				
	Person ID	Last name	First name	Date of birth
0	1	Smith	John	10/6/82
1	2	Williams	Bill	7/4/90
2	3	Williams	Jane	5/6/89
3	4	Jackson	Cathy	1/24/74
4	5	Johnson	Stuart	6/5/95

Dataset as Pandas DataFrame

	Person ID	Last name	First name	Date of birth
0	1	Smith	John	10/6/82
1	2	Williams	Bill	7/4/90
2	3	Williams	Jane	5/6/89
3	4	Jackson	Cathy	1/24/74
4	5	Johnson	Stuart	6/5/95
5	6	Smith	James	4/16/84
6	7	Anderson	Felicity	9/15/76
7	8	Christiansen	Liam	10/2/92
8	9	Carter	Nancy	2/5/86
9	10	Davidson	Christina	8/11/93

Work on the
whole
DataFrame
sequentially.



Dataset as Dask DataFrame

Partition 1 (Pandas DataFrame)

	Person ID	Last name	First name	Date of birth
0	1	Smith	John	10/6/82
1	2	Williams	Bill	7/4/90
2	3	Williams	Jane	5/6/89
3	4	Jackson	Cathy	1/24/74
4	5	Johnson	Stuart	6/5/95

Partition 2 (Pandas DataFrame)

	Person ID	Last name	First name	Date of birth
5	6	Smith	James	4/16/84
6	7	Anderson	Felicity	9/15/76
7	8	Christiansen	Liam	10/2/92
8	9	Carter	Nancy	2/5/86
9	10	Davidson	Christina	8/11/93

Data split
into
multiple
partitions
so the
work can
be shared

Work on
partition 1
in parallel.



Host 1

Work on
partition 2
in parallel.



Host 2

Dataset as Dask DataFrame

Partition 1 (Pandas DataFrame)

	Person ID	Last name	First name	Date of birth
0	1	Smith	John	10/6/82
1	2	Williams	Bill	7/4/90
2	3	Williams	Jane	5/6/89
3	4	Jackson	Cathy	1/24/74
4	5	Johnson	Stuart	6/5/95

Partition 2 (Pandas DataFrame)

	Person ID	Last name	First name	Date of birth
5	6	Smith	James	4/16/84
6	7	Anderson	Felicity	9/15/76
7	8	Christiansen	Liam	10/2/92
8	9	Carter	Nancy	2/5/86
9	10	Davidson	Christina	8/11/93

Data is split into multiple pieces.

Get partition 1

Get partition 2



Node 1

I'll work on partition 1.
Hey, node 2, you work
on partition 2.

Send result for
partition 1 to client



Client

Combined
result

OK, I'll work on
partition 2.

Send result for
partition 2 to client



Node 2

Worker nodes collaborate and
send their results back to the
client to be combined.



Server 1

Partition 1



Server 2

Partition 2

The Smiths span multiple partitions; either server 1 has to shuffle its Smith record to server 2 or vice versa.



Person ID	Last name	First name	Date of birth	
0	1	Smith	John	10/6/82
1	2	Williams	Bill	7/4/90
2	3	Williams	Jane	5/6/89
3	4	Jackson	Cathy	1/24/74
4	5	Johnson	Stuart	6/5/95

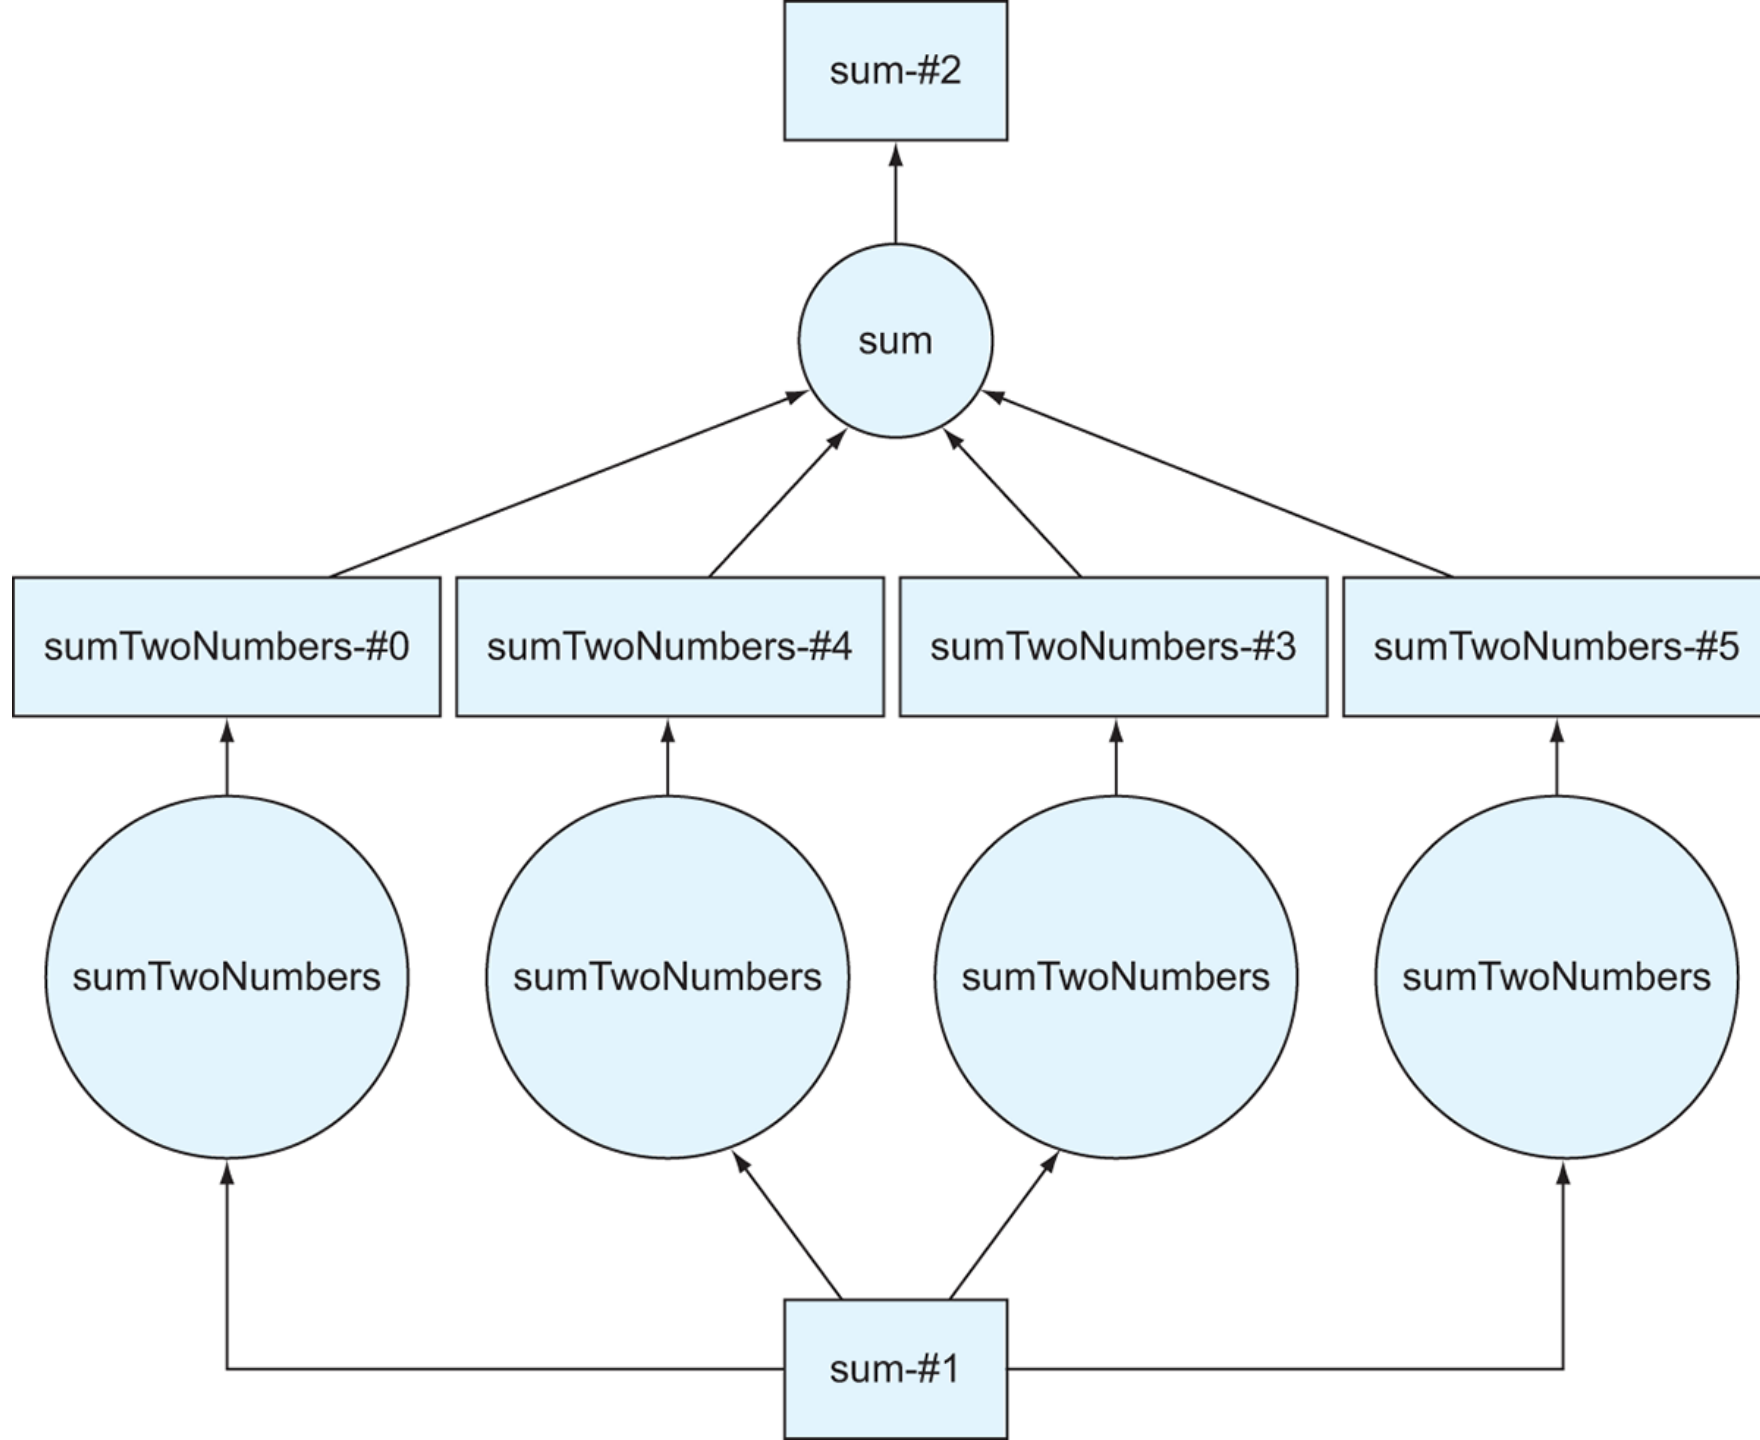
Person ID		Last name	First name	Date of birth
5	6	Smith	James	4/16/84
6	7	Anderson	Felicity	9/15/76
7	8	Christiansen	Liam	10/2/92
8	9	Carter	Nancy	2/5/86
9	10	Davidson	Christina	8/11/93

Both Williams are in the same partition!

The Index recycled back to 0 when the partition boundary was reached.

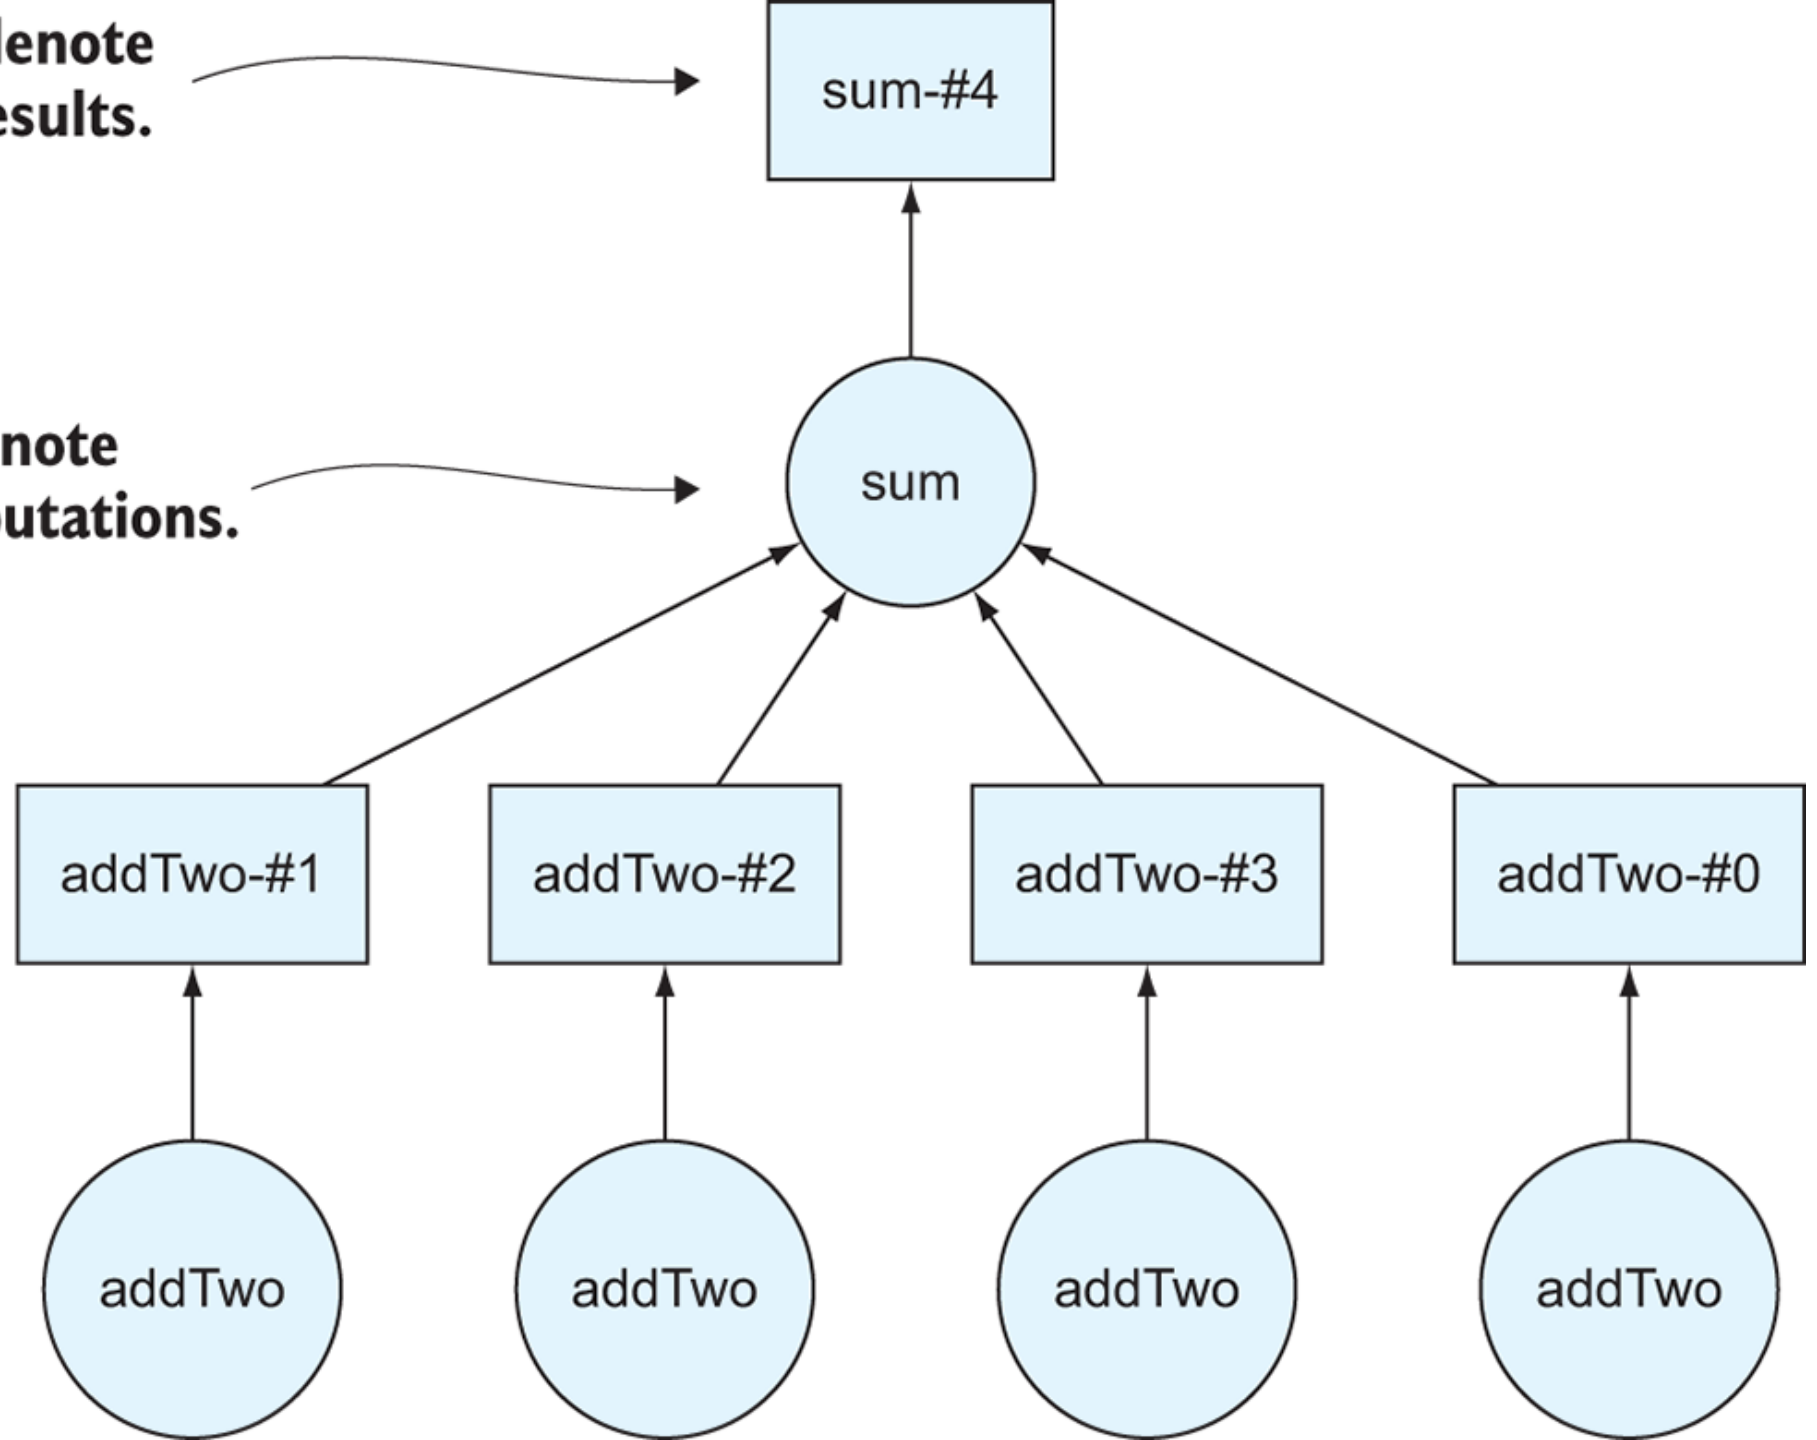


	Person ID	Last name	First name	Date of birth
0	1	Smith	John	10/6/82
1	2	Williams	Bill	7/4/90
2	3	Williams	Jane	5/6/89
3	4	Jackson	Cathy	1/24/74
4	5	Johnson	Stuart	6/5/95
0	6	Smith	James	4/16/84
1	7	Anderson	Felicity	9/15/76
2	8	Christiansen	Liam	10/2/92
3	9	Carter	Nancy	2/5/86
4	10	Davidson	Christina	8/11/93




**Square nodes denote
intermediate results.**


**Circle nodes denote
functions/computations.**



**All the data resides
on a single node.**





Worker node with data



Worker nodes without data

Other worker nodes that need data can easily saturate the single data node with requests, slowing our processing down to a halt.

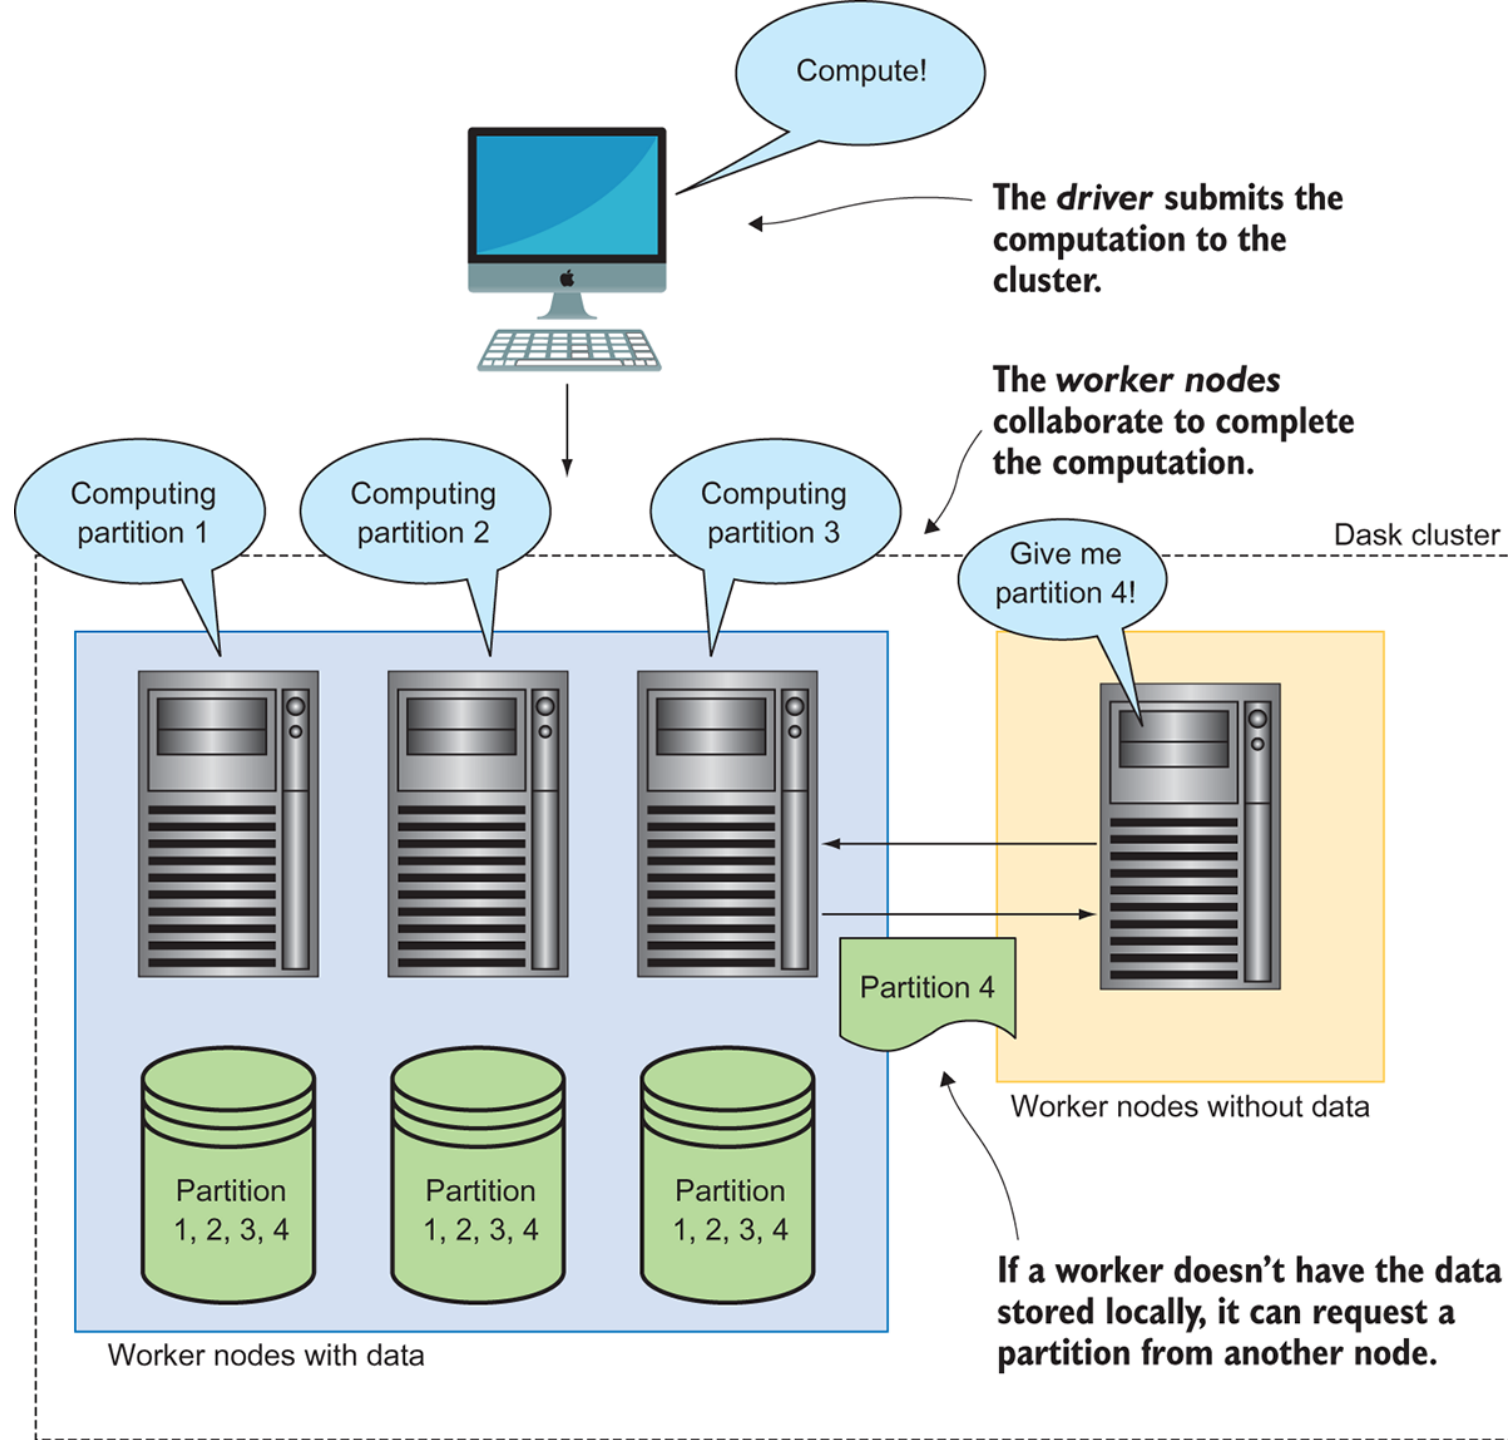
Worker nodes with data




Worker nodes without data

This worker node has three choices of nodes to get data from; it can choose the least busy node to request data, minimizing bottlenecks.

The data is stored in triplicate, eliminating single points of failure and performance bottlenecks.





```
import numpy as np
```

```
f = h5py.File('myfile.hdf5')  
x = np.array(f['/small-data'])  
x - x.mean(axis=1)
```

```
import dask.array as da
```

```
f = h5py.File('myfile.hdf5')  
x = da.from_array(f['/big-data'], chunks=(1000, 1000))  
x - x.mean(axis=1).compute()
```

Basic type	Available NumPy types	Comments
Boolean	bool	Elements are 1 byte in size.
Integer	int8, int16, int32, int64, int128, int	int defaults to the size of int in C for the platform.
Unsigned integer	uint8, uint16, uint32, uint64, uint128, uint	uint defaults to the size of unsigned int in C for the platform.
Float	float32, float64, float, longfloat	float is always a double-precision floating-point value (64 bits). longfloat represents large-precision floats. Its size is platform dependent.
Complex	complex64, complex128, complex	The real and complex elements of a complex64 are each represented by a single-precision (32-bit) value for a total size of 64 bits.
Strings	str, unicode	Unicode is always UTF32 (UCS4).
Object	object	Represents items in arrays as Python objects.
Records	void	Used for arbitrary data structures in record arrays.